

# The Processor

## COS375 / ELE375

David I. August

Acknowledgements:  
Mohammad Shahrad  
Margaret Martonosi



# Fall 2022



Prof. David August ([august@princeton.edu](mailto:august@princeton.edu))  
Office Hours: After class and by appointment

# Course Staff: Your Other Instructor

1999-Present:

- Professor of Computer Science at Princeton
- Compiler and Computer Architecture Research (more later)
- Lead Liberty Research Group



(at age 15)

1993-1999:

- M.S. and Ph.D. Electrical Engineering from University of Illinois
- Ph.D. Thesis Topic: Predication (will discuss briefly later in semester)

Ongoing consulting and research relationships with companies like Intel, Apple, Google, Samsung, Sony, etc.



# GodBolt.Org

COMPILER EXPLORER

Add... More Templates

C source #1

```
A -> C + -> V C C
1 int FromLecture3(int h, int *A) {
2     A[12] = h + A[8];
3     return A[12];
4 }
```

Not Optimized

RISC-V (64-bits) gcc 15.2.0 (Editor #1)

RISC-V (64-bits) gcc 15.2.0 (Editor #1) Compiler options... -O2

A -> C + -> V C C

```
FromLecture3:
1 lw    a5,32(a1)
2 addw a0,a5,a0
3 sw    a0,48(a1)
4 ret
```

Optimized

RISC-V (64-bits) gcc 15.2.0 (Editor #1)

RISC-V (64-bits) gcc 15.2.0 (Editor #1) Compiler options... -O2

A -> C + -> V C C

```
FromLecture3:
1 addi sp,sp,-32
2 sd   ra,24(sp)
3 sd   s0,16(sp)
4 addi s0,sp,32
5 mv   a5,a0
6 sd   a1,-32(s0)
7 sw   a5,-20(s0)
8 ld   a5,-32(s0)
9 addi a5,a5,32
10 lw   a4,0(a5)
11 ld   a5,-32(s0)
12 addi a5,a5,48
13 lw   a3,-20(s0)
14 addw a4,a3,a4
15 sext.w a4,a4
16 sw   a4,0(a5)
17 ld   a5,-32(s0)
18 lw   a5,48(a5)
19 mv   a0,a5
20 ld   ra,24(sp)
21 ld   s0,16(sp)
22 addi sp,sp,32
23 jr   ra
```

# Let's Build a RISC-V Processor!

Support a subset of instructions:

- Arithmetic/logical (R-type)
- Memory: lw, sw
- Control flow: beq

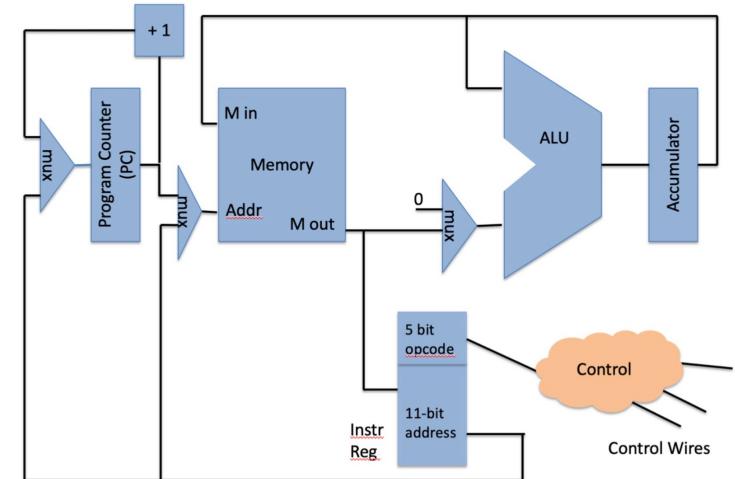
Datapath, then control



# Datapath

## Datapath

1949 Datapath:



“The component of the processor that performs arithmetic operations” – P&H p. 20

**datapath** The component of the processor that performs arithmetic operations.

## Datapath

“The collection of state elements, computation elements, and interconnections that transform data in the processor during execution.” – Me

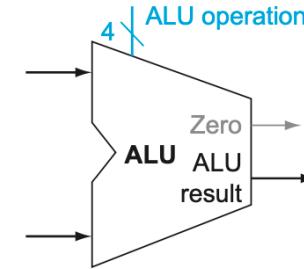


# Datapath Computational Elements

Blue vs. Black?

## ALU:

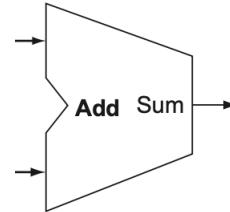
- Supports different operations
- In RISC-V, ALU has an output for zero equality. Why?



| ALU control lines | Function |
|-------------------|----------|
| 0000              | AND      |
| 0001              | OR       |
| 0010              | add      |
| 0110              | subtract |

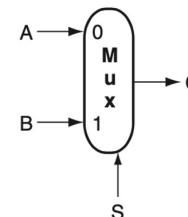
## Adder:

- Not an ALU, just add
- Why useful in RISC-V?



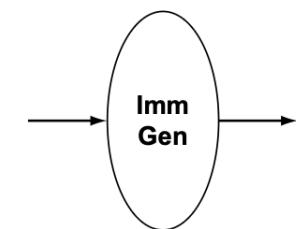
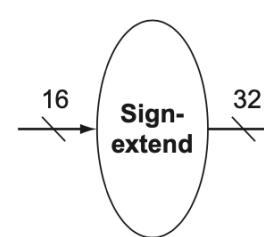
## The Magical Mux:

- Mux selects an input based on the control line.
- Why useful in any datapath?



## Sign-Extender and Imm-Gen:

- This sign-extender replicates bit [15] 16 times.
- What does Imm Gen do in RISC-V?



| I  | imm[11:0]    | rs1 | funct3 | rd     | Opcode             |
|----|--------------|-----|--------|--------|--------------------|
| S  | imm[11:5]    | rs2 | rs1    | funct3 | imm[4:0] opcode    |
| SB | imm[12:10:5] | rs2 | rs1    | funct3 | imm[4:1 11] opcode |



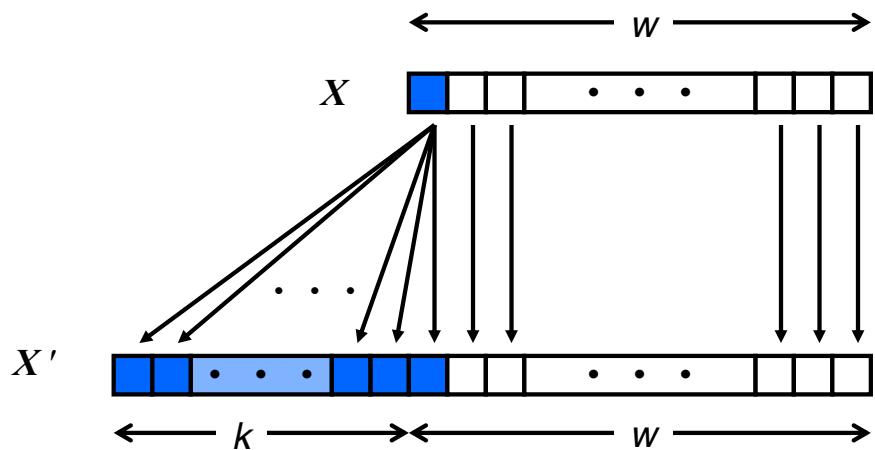
ALL COMBINATIONAL

# Sign Extension vs. Zero Extension

```
char minusFour = -4;  
short moreBits;  
moreBits = (short)minusFour;
```

Given w bit signed integer,  
return equivalent w+k bit signed integer.

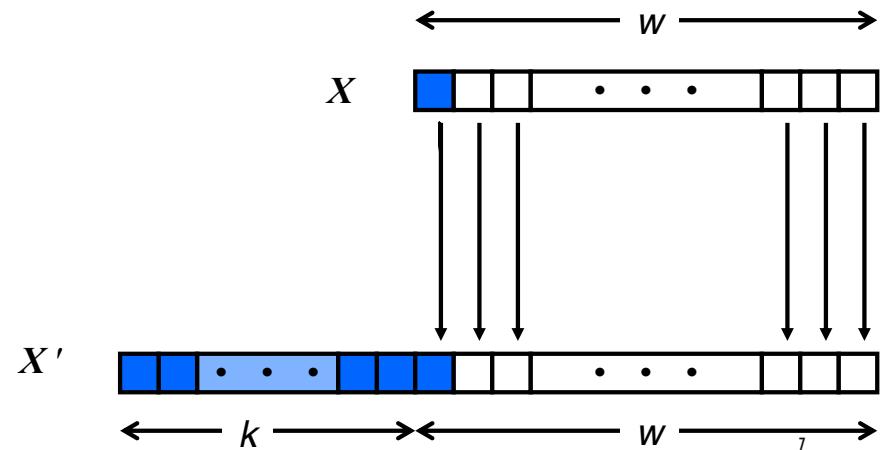
Sign Extend:



```
unsigned char Four = 4;  
unsigned short moreBits;  
moreBits = (unsigned short)Four;
```

Given w bit unsigned integer,  
return equivalent w+k bit unsigned integer.

Sign Extend:



# REVIEW SLIDES



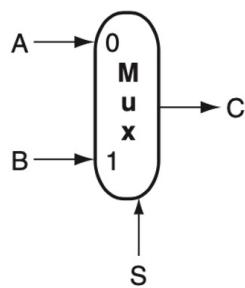
# Combinational Circuits

- The ALU, Adder, Sign-Extender, and Mux are all Combinational Circuits.
- A Combinational Circuit is a digital circuit whose outputs depend solely on the present combination of the circuit input's values. No memory (i.e., no state).
- A combinational circuit implements a Boolean function.
- A Boolean function can be represented by enumerating the function output value for all possible input values: A Truth Table

Mux Example:

Inputs:  $s$ ,  $i_0$ ,  $i_1$

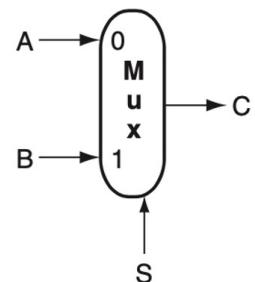
Outputs:  $out$



| S | A | B | C |
|---|---|---|---|
| 0 | 0 | 0 |   |
| 0 | 0 | 1 |   |
| 0 | 1 | 0 |   |
| 0 | 1 | 1 |   |
| 1 | 0 | 0 |   |
| 1 | 0 | 1 |   |
| 1 | 1 | 0 |   |
| 1 | 1 | 1 |   |

| S | A | B | C |
|---|---|---|---|
|   |   |   |   |
|   |   |   |   |
|   |   |   |   |
|   |   |   |   |
|   |   |   |   |
|   |   |   |   |
|   |   |   |   |

# Computation Element: Mux Boolean Expression



- A Boolean function can be represented by enumerating the function output value for all possible input values: A Truth Table
- A Boolean function can be represented by Boolean Algebra.
- Truth Table to Sum-Of-Products Boolean Expression Conversion:

| S | A | B | C |
|---|---|---|---|
| 0 | 0 | 0 | 0 |
| 0 | 0 | 1 | 0 |
| 0 | 1 | 0 | 1 |
| 0 | 1 | 1 | 1 |
| 1 | 0 | 0 | 0 |
| 1 | 0 | 1 | 1 |
| 1 | 1 | 0 | 0 |
| 1 | 1 | 1 | 1 |

$S'AB'$

$S'AB$

$SA'B$

$SAB$

$S'AB' + S'AB + SA'B + SAB$

| S | A | B | C | $S'A'$ | $S'A$ | $SB'$ | $SB$ | $S'A+SB$ |
|---|---|---|---|--------|-------|-------|------|----------|
| 0 | 0 | X | 0 | 1      | 0     | 0     | 0    | 0        |
| 0 | 1 | X | 1 | 0      | 1     | 0     | 0    | 1        |
| 1 | X | 0 | 0 | 0      | 0     | 1     | 0    | 0        |
| 1 | X | 1 | 1 | 0      | 0     | 0     | 1    | 1        |

$S'A+SB$

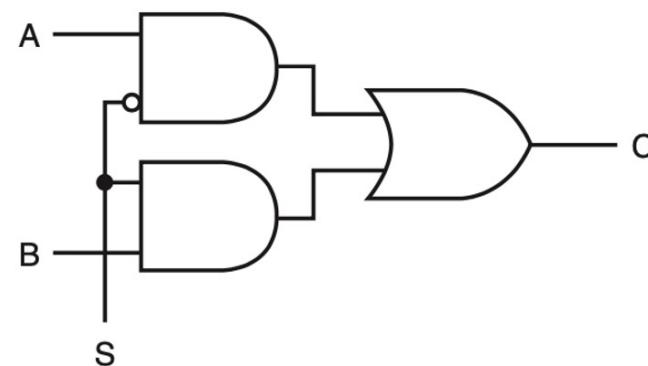
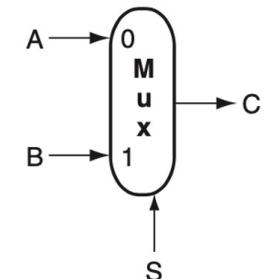


# Computation Element: Mux Circuit!

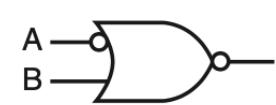
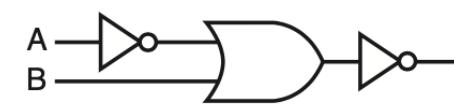
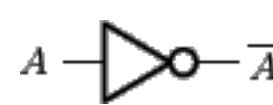
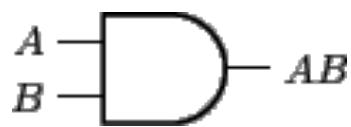
- A Boolean function can be represented by enumerating the function output value for all possible input values: A Truth Table
- A Boolean function can be represented by Boolean Algebra.
- A Boolean function can be represented and implemented by a Circuit.

| S | A | B | C |
|---|---|---|---|
| 0 | 0 | X | 0 |
| 0 | 1 | X | 1 |
| 1 | X | 0 | 0 |
| 1 | X | 1 | 1 |

$$C = S'A + SB$$



Symbology:



THIS IS AS LOW AS WE GO...

# You can physically build these circuits!

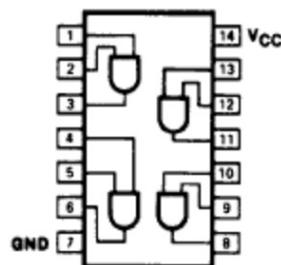
Dedicated Integrated Circuits (IC)

7408 IC

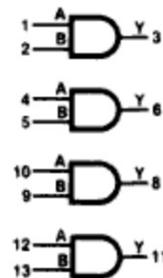


7408 • 74S08 • 74LS08  
QUAD 2-INPUT AND GATE

PIN ASSIGNMENT



LOGIC DIAGRAM

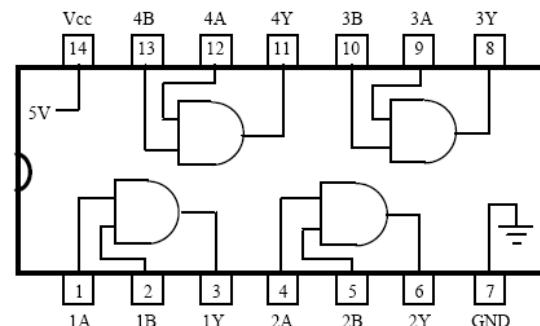


TRUTH TABLE

| INPUTS |   | OUTPUT |
|--------|---|--------|
| A      | B | Y      |
| L      | L | L      |
| L      | H | L      |
| H      | L | L      |
| H      | H | H      |

H = HIGH voltage level  
L = LOW voltage level

Pinout of the 7408 IC

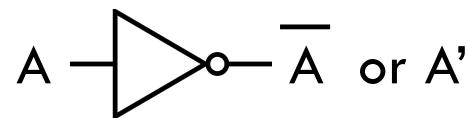


# Boolean Algebra

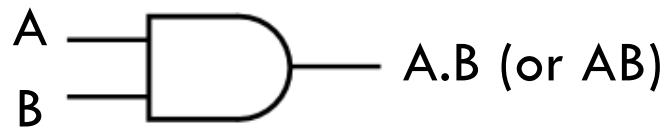
- Introduced by George Boole in his 1847 book
- Fundamental operators:
- NOT (negation)  $[0 \rightarrow 1, 1 \rightarrow 0]$



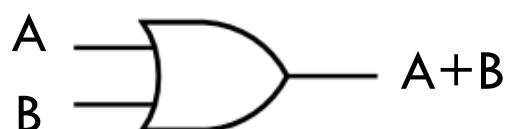
George Boole



- AND (output is 1 if both inputs are 1)



- OR (output is 1 if any input is 1)



These three operators are Boolean Functions.

These three simple Boolean Functions are universal (e.g., they can build all other functions).



# Boolean Algebra Laws

MUX:

$$S'AB' + S'AB + SA'B + SAB = S'A + SB$$

$$S'AB' + S'AB + SA'B + SAB = S'A(B' + B) + S(A' + A)B \quad (\text{Distribution})$$

$$S'A(B' + B) + S(A' + A)B = S'A + SB \quad (\text{Inversion})$$

Convince yourself these are true with truth tables:  $\overline{A + B} = \overline{A} \cdot \overline{B}$

Identity:  $A \cdot 1 = A, A + 0 = A$

Zero and one:  $A \cdot 0 = 0, A + 1 = 1$

Inversion:  $A \cdot \overline{A} = 0, A + \overline{A} = 1$

Idempotence:  $A \cdot A = A, A + A = A$

| A | B | $A + B$ | $\overline{A + B}$ | $\overline{A}$ | $\overline{B}$ | $\overline{A} \cdot \overline{B}$ |
|---|---|---------|--------------------|----------------|----------------|-----------------------------------|
| 0 | 0 | 0       | 1                  | 1              | 1              | 1                                 |
| 0 | 1 | 1       | 0                  | 1              | 0              | 0                                 |
| 1 | 0 | 1       | 0                  | 0              | 1              | 0                                 |
| 1 | 1 | 1       | 0                  | 0              | 0              | 0                                 |

Commutativity:  $A \cdot B = B \cdot A, A + B = B + A$

Associativity:  $A \cdot (B \cdot C) = (A \cdot B) \cdot C, A + (B + C) = (A + B) + C$

Distribution:  $A \cdot (B + C) = A \cdot B + A \cdot C, A + (B \cdot C) = (A + B) \cdot (A + C)$

DeMorgan's theorems:  $\overline{A \cdot B} = \overline{A} + \overline{B}, \overline{A + B} = \overline{A} \cdot \overline{B}$

(You'll find this one very useful.)



# END REVIEW SLIDES

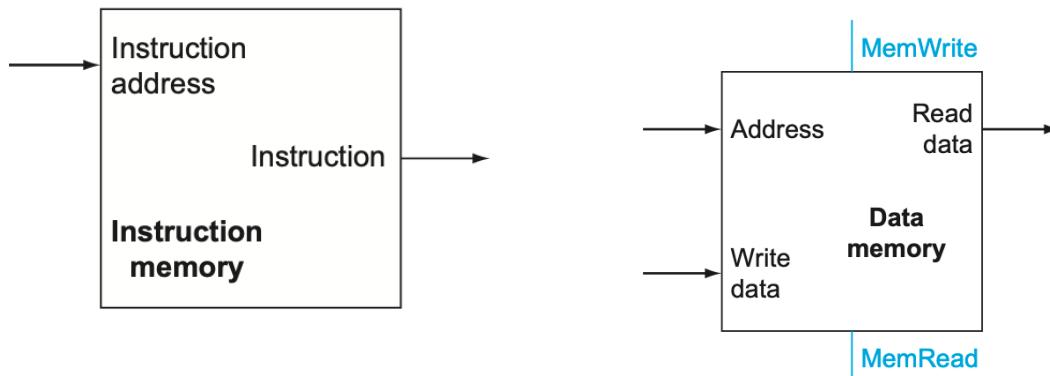
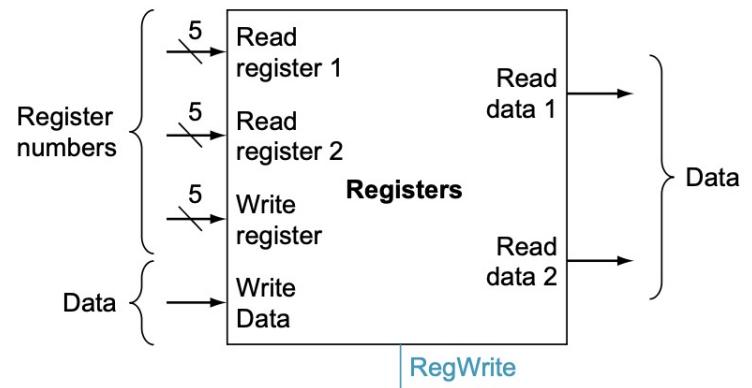


# Datapath State Elements

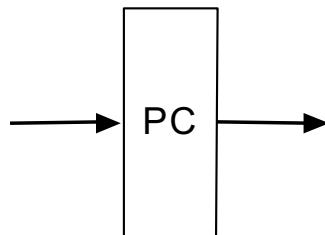
## RISC-V Register File:

- Why “Read register 1”, “Read register 2” and “Write register” inputs?
- Why 5 bits each?

## Instruction and Data Memory (for now and again much later):

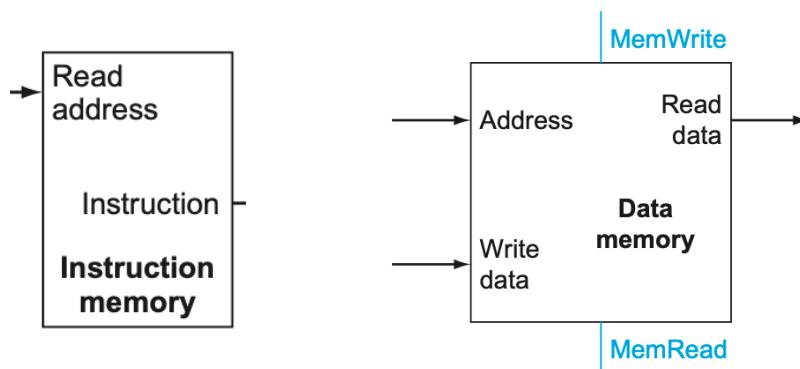
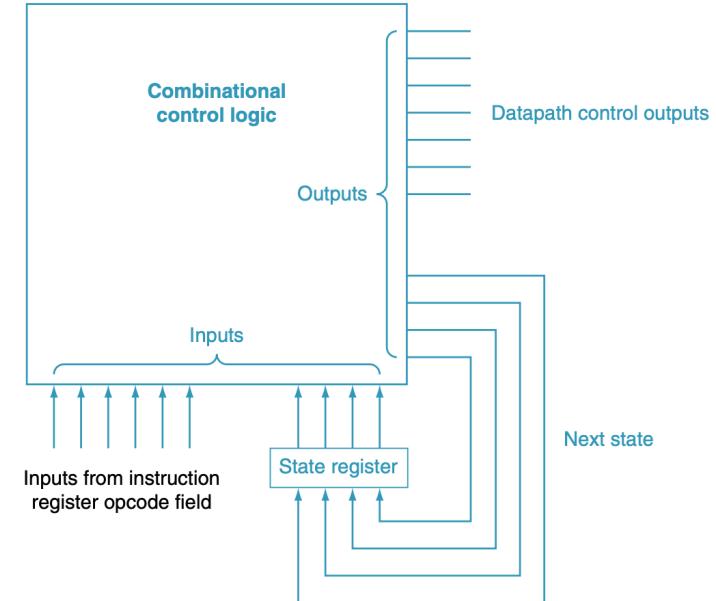


## Program Counter (PC):



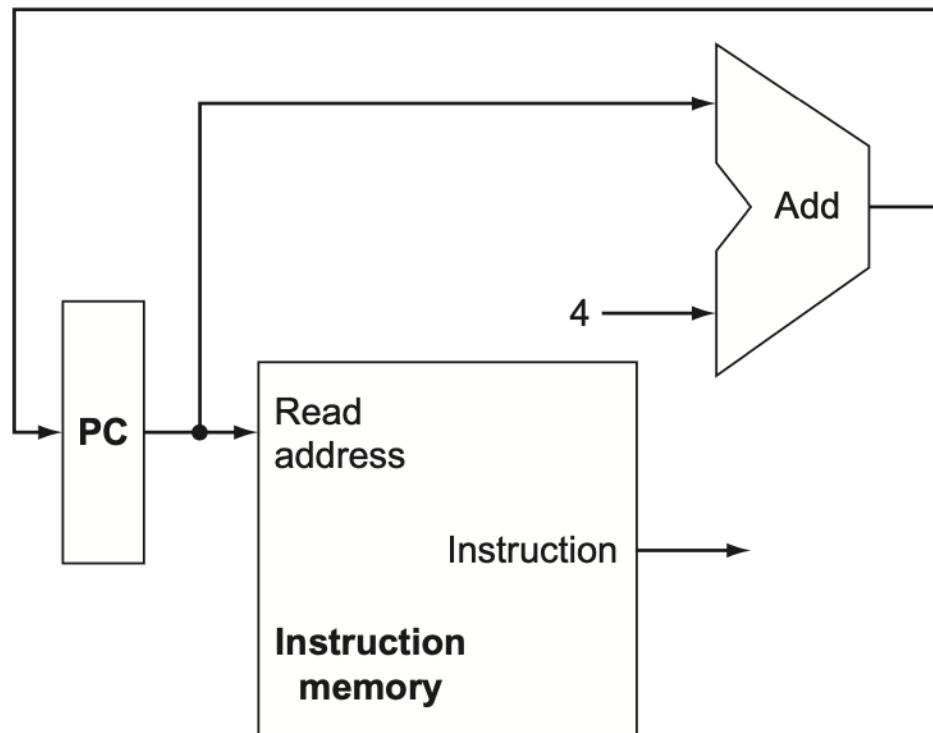
# REVIEW: State Elements

- State Elements are Sequential Circuits, not Combinational Circuits.
- Sequential Circuit output is a function of state (and often also input).
- How does one design and build Sequential Circuits?
- We will see when we build ADVANCED Control. For now, let's build the Datapath...



Let's start building our processor, datapath first!

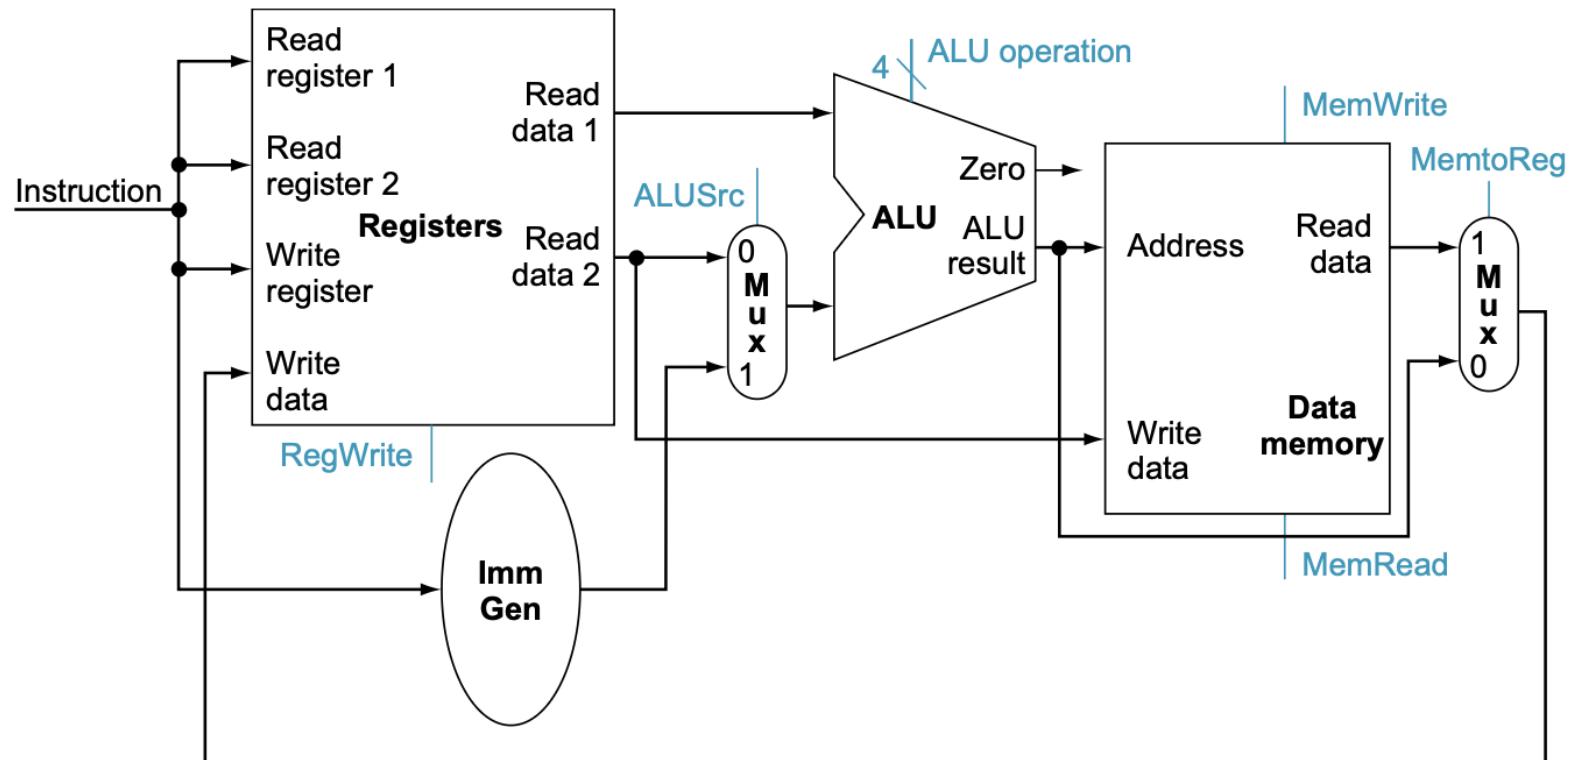
# Fetching Instructions (no branching)



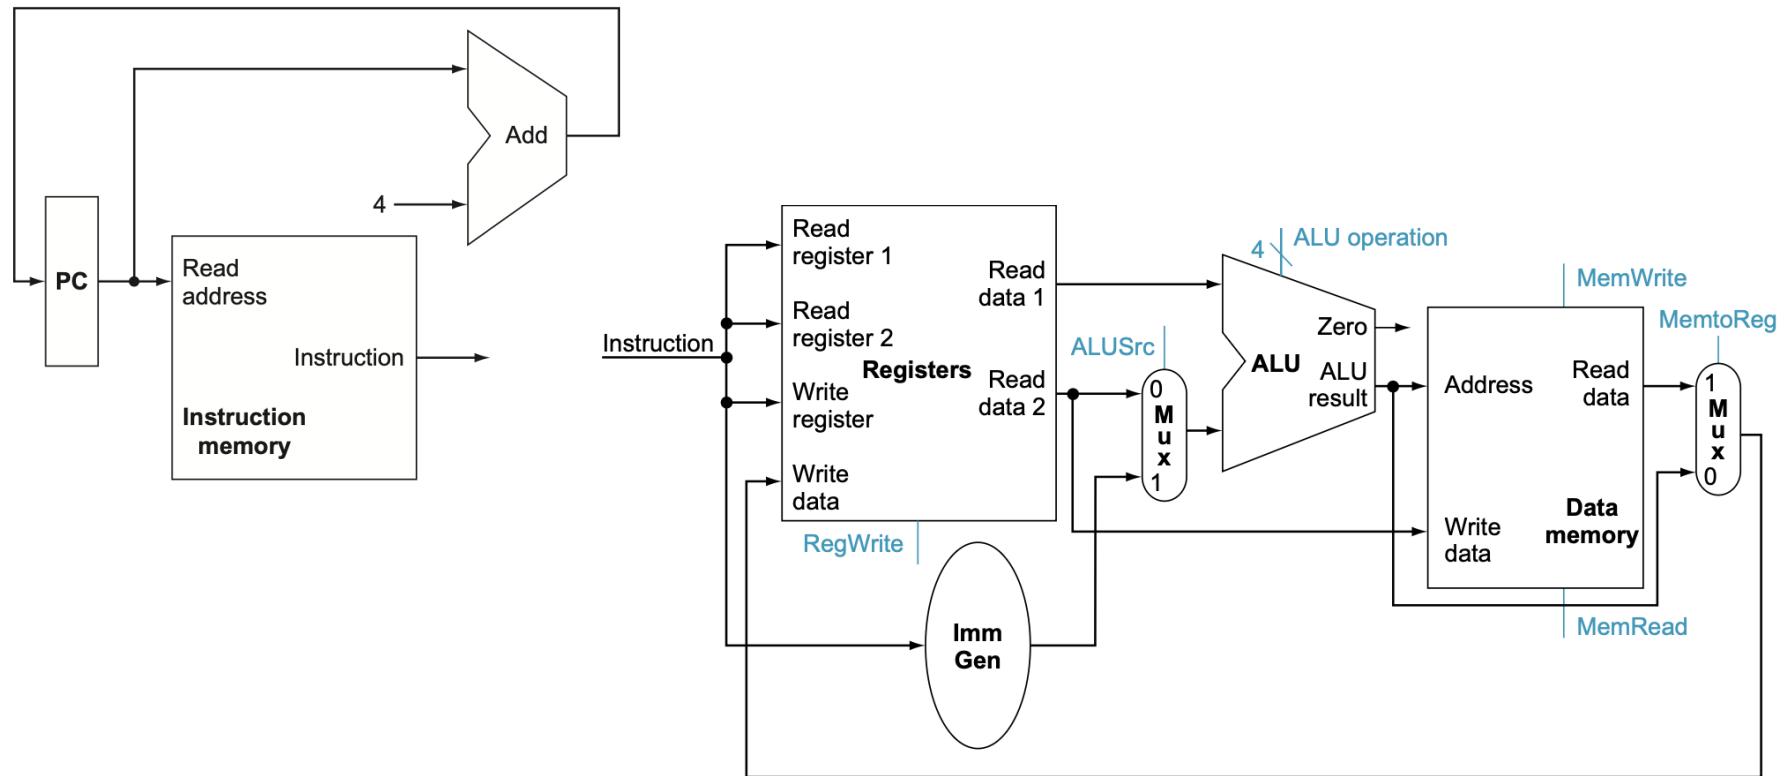
This part looks like EDSAC fetch



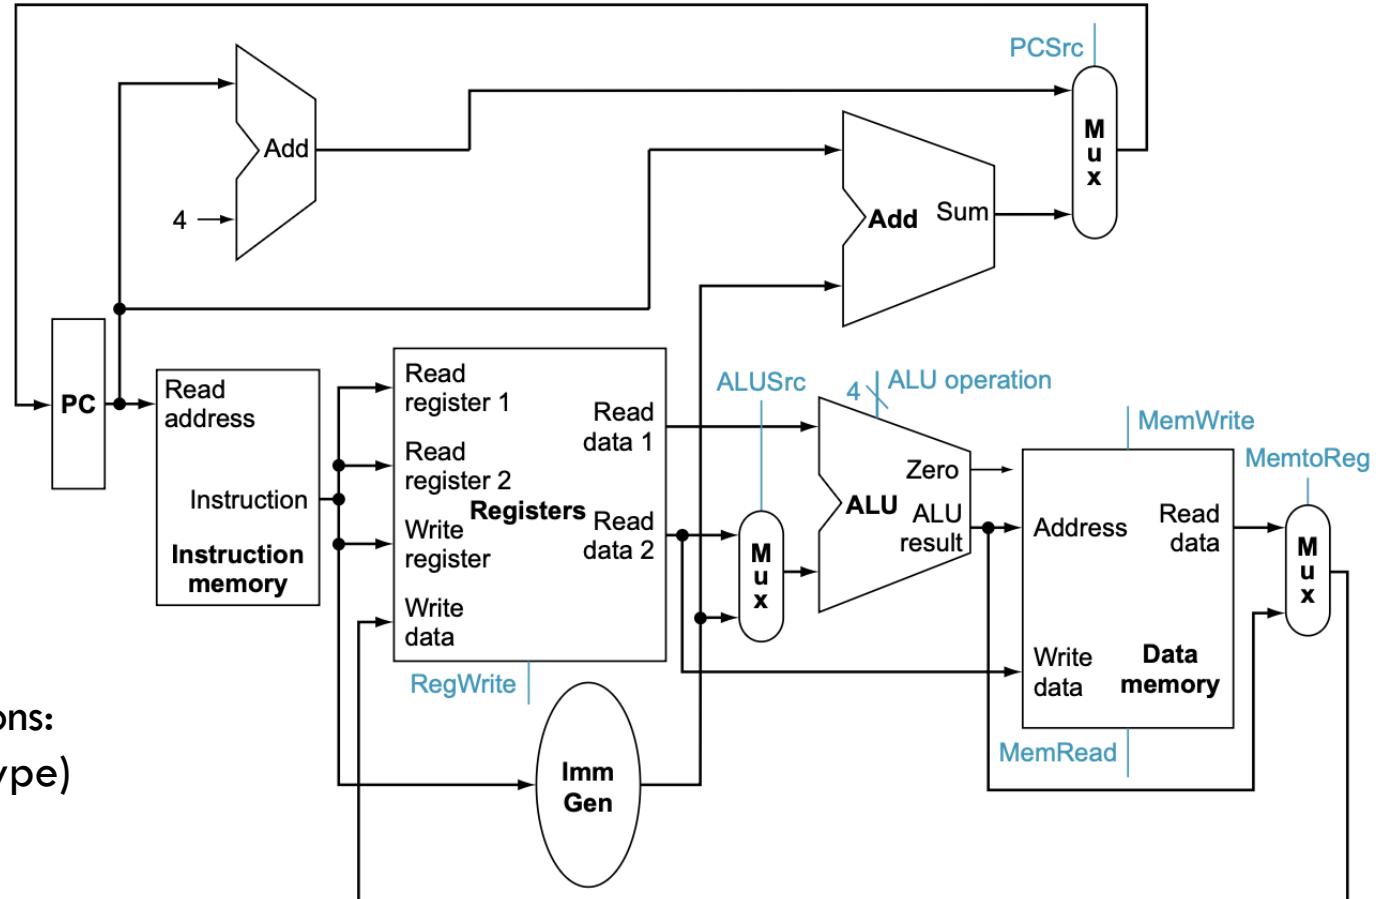
# Memory and R-Type Datapath



# Bring Fetch and Mem/R-Type Datapaths Together:



# Adding Instruction Fetch With Branches

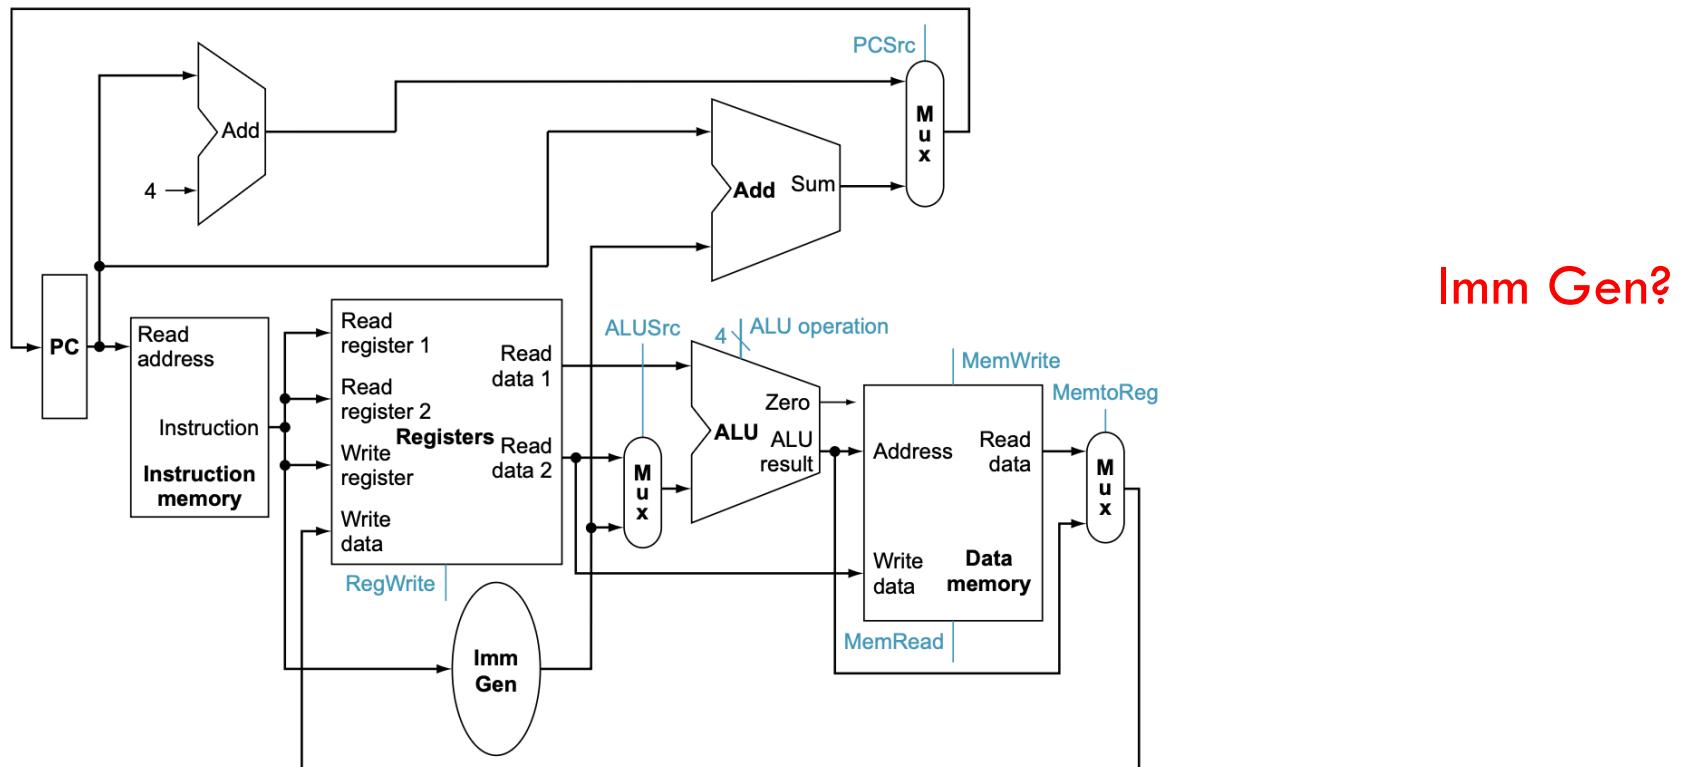


Supports a subset of instructions:

- Arithmetic/logical (R-type)
- Memory: `lw`, `sw`
- Control flow: `beq`

# Adding Instruction Fetch With Branches

|         | 31        | 27 | 26 | 25           | 24 | 20  | 19 | 15  | 14 | 12     | 11 | 7           | 6 | 0      |
|---------|-----------|----|----|--------------|----|-----|----|-----|----|--------|----|-------------|---|--------|
| R-type: | <b>R</b>  |    |    | funct7       |    | rs2 |    | rs1 |    | funct3 |    | rd          |   | Opcode |
| lw:     | <b>I</b>  |    |    | imm[11:0]    |    |     |    | rs1 |    | funct3 |    | rd          |   | Opcode |
| sw:     | <b>S</b>  |    |    | imm[11:5]    |    | rs2 |    | rs1 |    | funct3 |    | imm[4:0]    |   | opcode |
| beq:    | <b>SB</b> |    |    | imm[12 10:5] |    | rs2 |    | rs1 |    | funct3 |    | imm[4:1 11] |   | opcode |



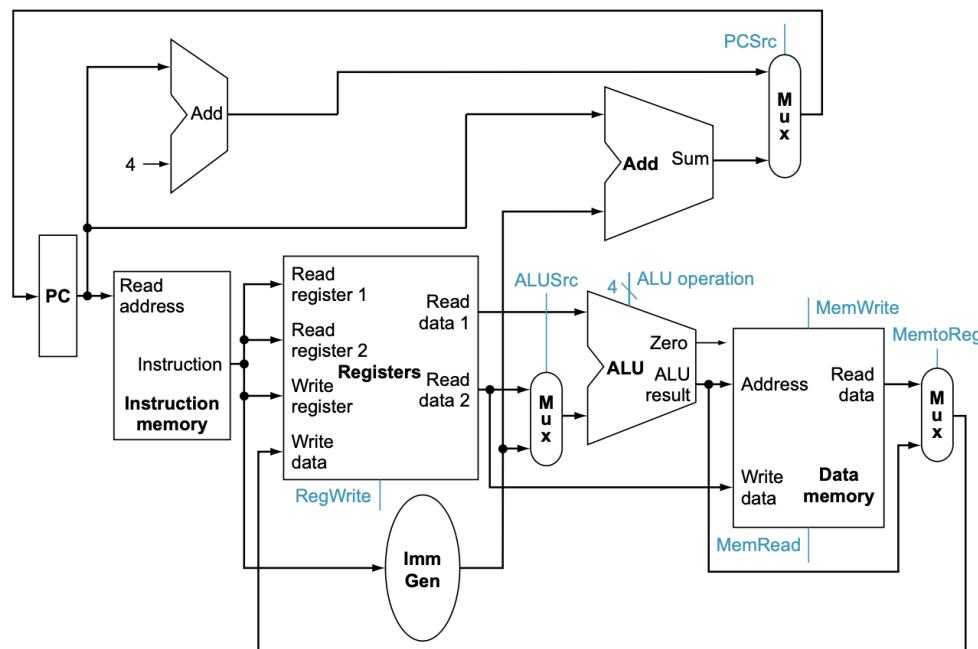
# Adding Instruction Fetch With Branches

|           | 31 | 27                    | 26 | 25  | 24 | 20  | 19 | 15     | 14 | 12          | 11 | 7      | 6 | 0 |
|-----------|----|-----------------------|----|-----|----|-----|----|--------|----|-------------|----|--------|---|---|
| R-type: R |    | funct7                |    | rs2 |    | rs1 |    | funct3 |    | rd          |    | Opcode |   |   |
| lw: I     |    | imm[11:0]             |    |     |    | rs1 |    | funct3 |    | rd          |    | Opcode |   |   |
| sw: S     |    | imm[11:5]             |    | rs2 |    | rs1 |    | funct3 |    | imm[4:0]    |    | opcode |   |   |
| beq: SB   |    | imm[12 10:5]          |    | rs2 |    | rs1 |    | funct3 |    | imm[4:1 11] |    | opcode |   |   |
| jal: UJ   |    | imm[20 10:1 11 19:12] |    |     |    |     |    |        |    | rd          |    | opcode |   |   |

jal

UJ Jump & Link

$R[rd] = PC+4$ ;  $PC = PC + \{imm, 1b'0\}$



Can this datapath support jal?

# “Architecture” Coined (1964)

• The term *architecture* is used here to describe the attributes of a system as seen by the programmer, i.e., the conceptual structure and functional behavior, as distinct from the organization of the data flow and controls, the logical design, and the physical implementation.

## Architecture or Instruction Set Architecture (ISA):

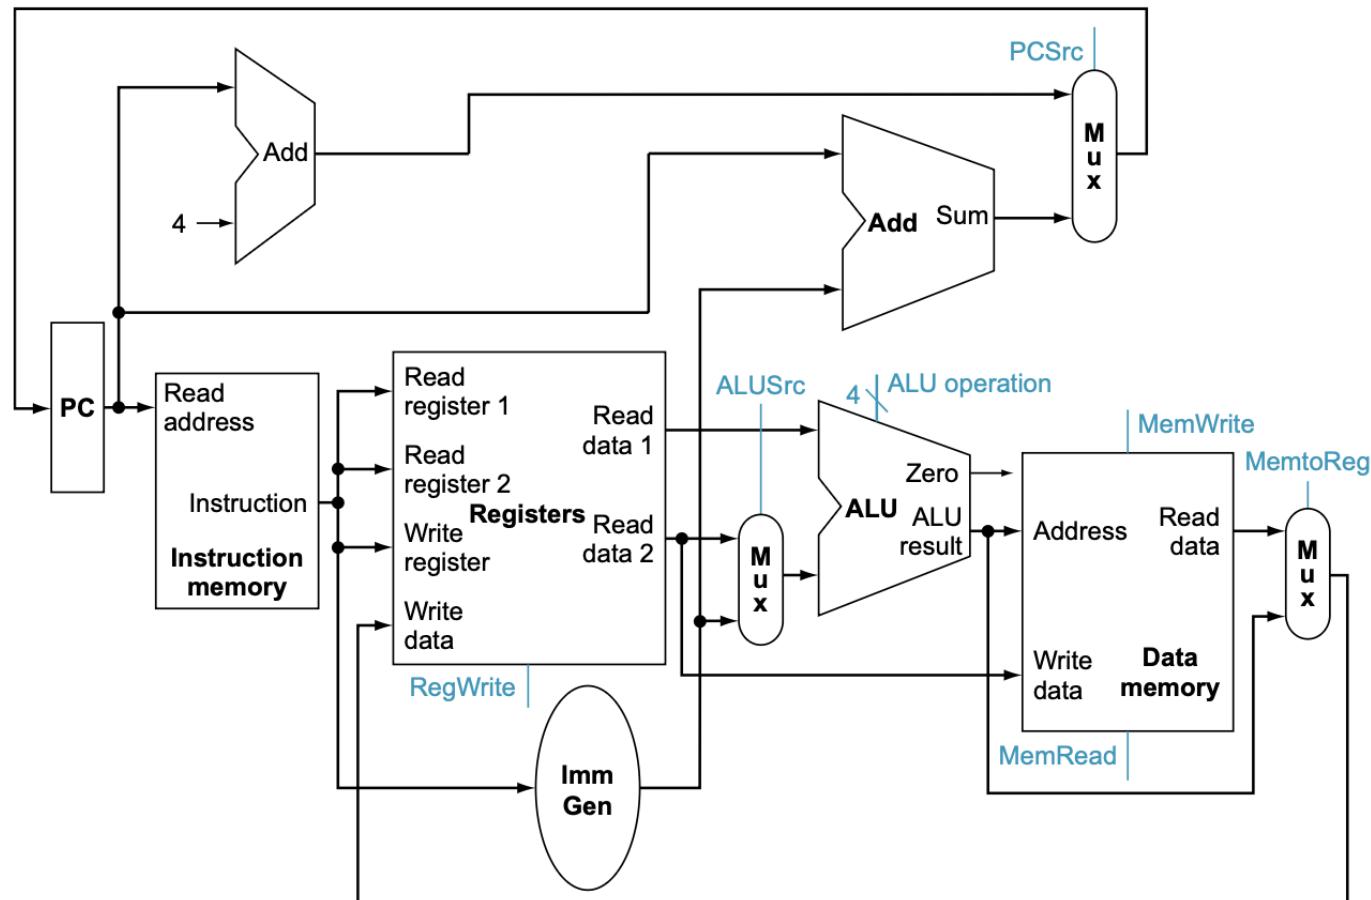
- A “contract” between HW and SW
- Examples: ARM, x86, MIPS, RISC-V, PowerPC, ...

## Implementation or Microarchitecture:

- A physical instance of an architecture
- Apple M1 vs. M2, Intel Core i5-1038NG7 vs Xeon ES-2620



# Full Datapath: Architectural/Microarchitectural State



# MIPS Instruction Quirk

What is MIPS?

| RISC-V: | R-type: R | 31     | 27        | 26  | 25  | 24  | 20     | 19       | 15 | 14 | 12 | 11 | 7 | 6 | 0      |
|---------|-----------|--------|-----------|-----|-----|-----|--------|----------|----|----|----|----|---|---|--------|
|         | lw: I     | funct7 |           | rs2 | rs1 |     | funct3 |          | rd |    |    |    |   |   | Opcode |
|         | sw: S     |        | imm[11:0] |     | rs1 |     | funct3 |          | rd |    |    |    |   |   | Opcode |
|         |           |        | imm[11:5] |     | rs2 | rs1 | funct3 | imm[4:0] |    |    |    |    |   |   | opcode |

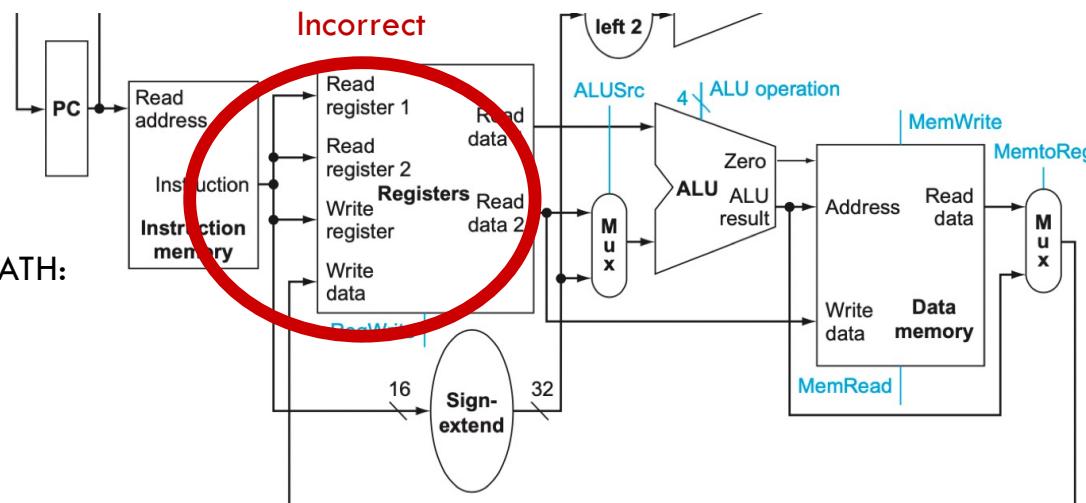
In MIPS, the Destination Register may be in different locations:

- All R-Types use rd (bits 11-15)
- Loads use rt (bits 16-20)

|            |        |        |        |        |        |       |
|------------|--------|--------|--------|--------|--------|-------|
| R-type: R: | 6 bits | 5 bits | 5 bits | 5 bits | 6 bits |       |
|            | op     | rs     | rt     | rd     | shamt  | funct |

|               |        |        |        |                     |
|---------------|--------|--------|--------|---------------------|
| Iw and sw: I: | 6 bits | 5 bits | 5 bits | 6 bits              |
|               | op     | rs     | rt     | address / immediate |

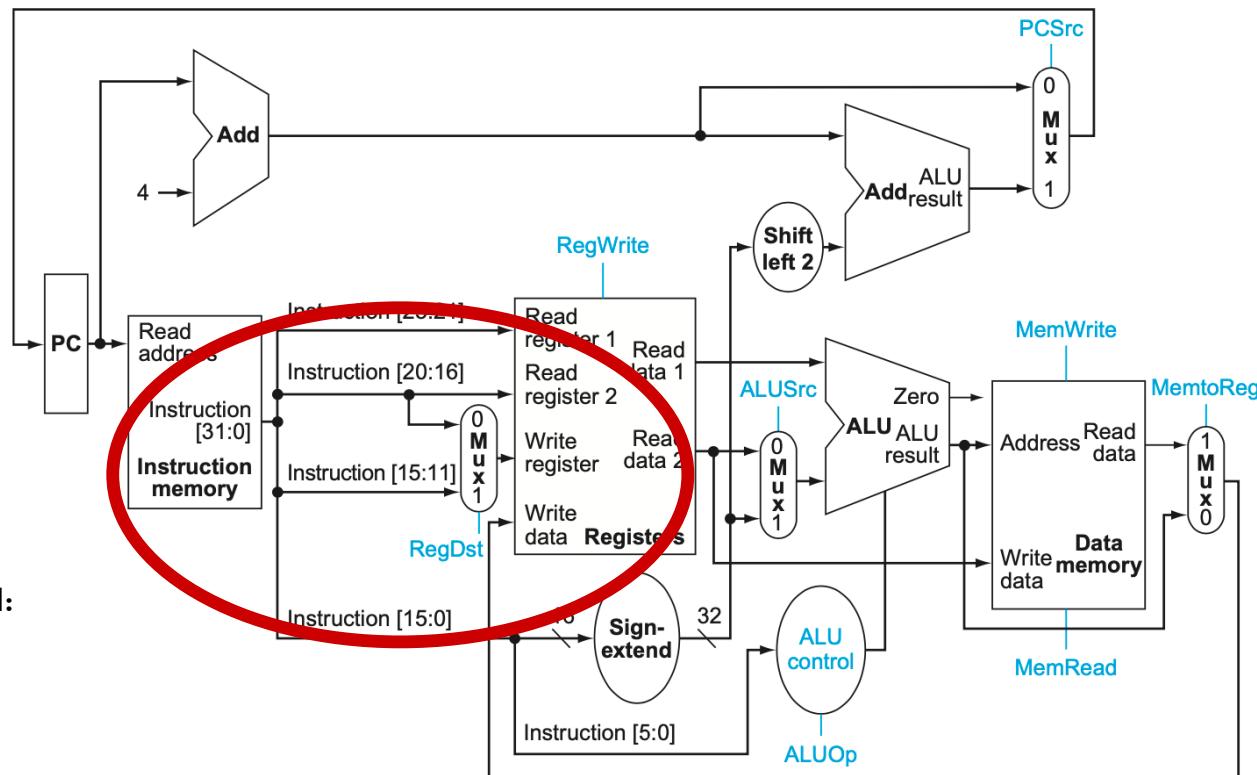
BROKEN MIPS DATAPATH:



# The magic of the Mux!

MIPS:

|    | 6 bits | 5 bits | 5 bits | 5 bits              | 5 bits | 6 bits |
|----|--------|--------|--------|---------------------|--------|--------|
| R: | op     | rs     | rt     | rd                  | shamt  | funct  |
| I: | op     | rs     | rt     | address / immediate |        |        |



FIXED MIPS DATAPATH:

ICQ: Why do we use different ISAs rather than have a monolithic version?

# Control

**control** The component of the processor that commands the datapath, memory, and I/O devices according to the instructions of the program.

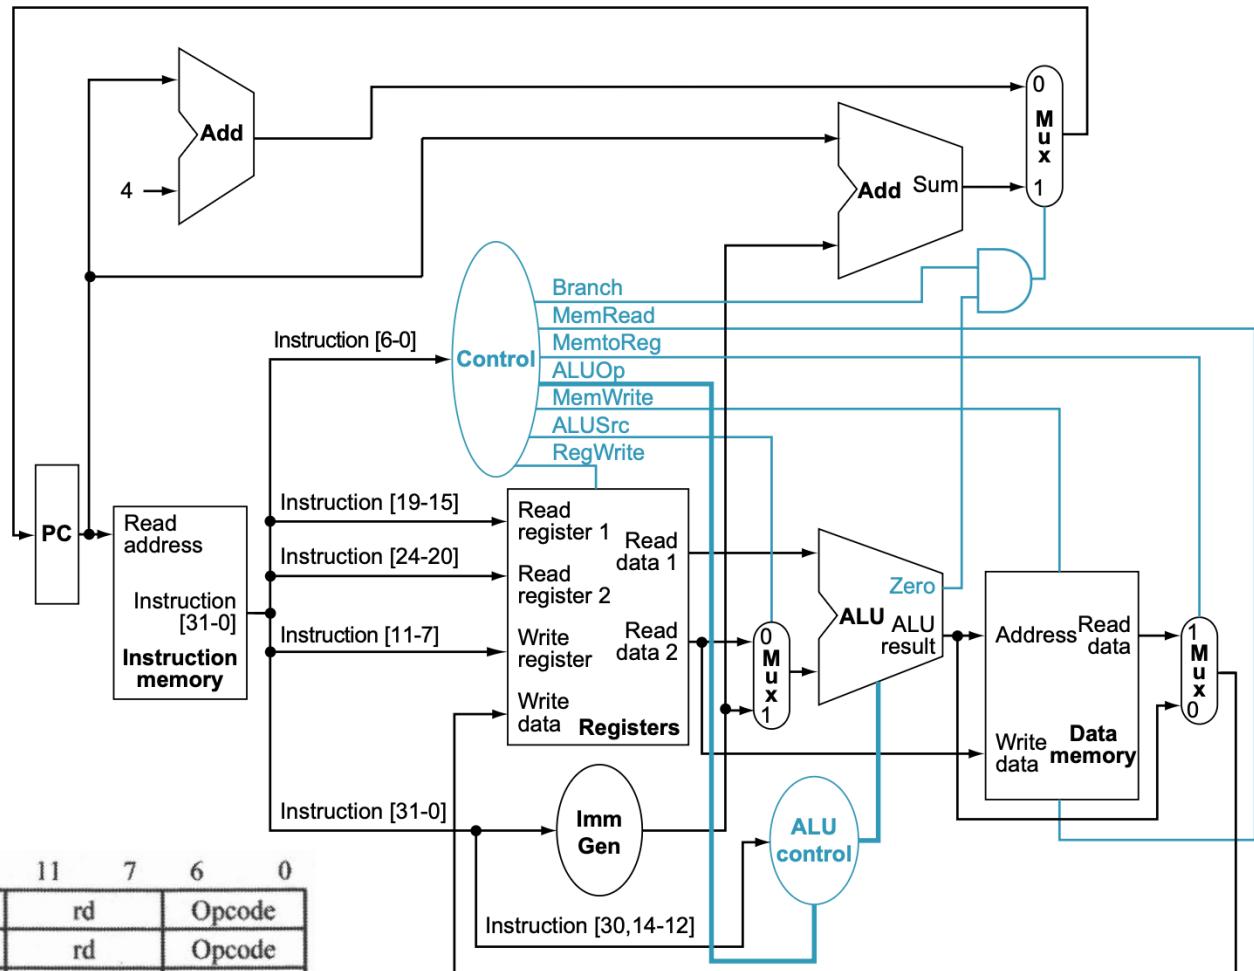
“The component of the processor that commands the datapath, memory, and I/O devices according to the instructions of the program.” – P&H p. 21

I prefer: “The component of the processor that commands the datapath according to the instructions of the program.”



# Full RISC-V Datapath with Control

- Supports a subset of instructions:
  - Arithmetic/logical (R-type)
  - Memory: lw, sw
  - Control flow: beq
- Control generates signals: Branch, MemRead, MemtoReg, ALUOp, MemWrite, ALUSrc, RegWrite
- Control uses opcode: Instruction[6-0]



| 31 | 27           | 26 | 25 | 24  | 20 | 19  | 15 | 14     | 12          | 11 | 7      | 6      | 0 |
|----|--------------|----|----|-----|----|-----|----|--------|-------------|----|--------|--------|---|
| R  | funct7       |    |    | rs2 |    | rs1 |    | funct3 |             | rd |        | Opcode |   |
| I  | imm[11:0]    |    |    |     |    | rs1 |    | funct3 |             | rd |        | Opcode |   |
| S  | imm[11:5]    |    |    | rs2 |    | rs1 |    | funct3 | imm[4:0]    |    | opcode |        |   |
| SB | imm[12:10:5] |    |    | rs2 |    | rs1 |    | funct3 | imm[4:1 11] |    | opcode |        |   |

Remember: THIS IS A LITTLE OFF.

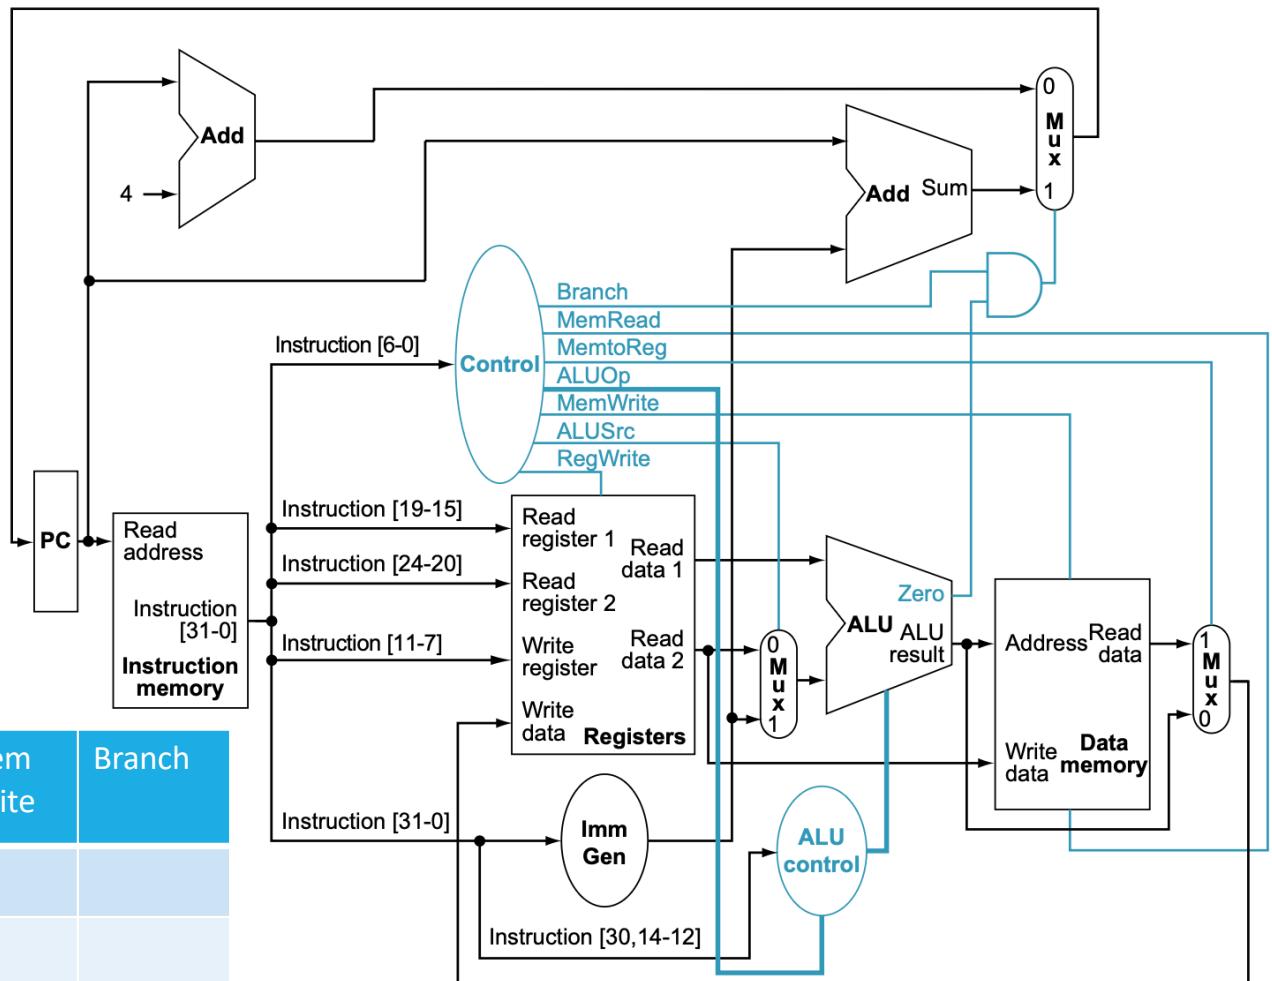
# Control Signals in Full Datapath

| Signal name | Effect when deasserted   | Effect when asserted  |
|-------------|--|---|
| RegWrite    | None.  | The register on the Write register input is written with the value on the Write data input.             |
| ALUSrc      | The second ALU operand comes from the second register file output (Read data 2). | The second ALU operand is the sign-extended, 12 bits of the instruction.                                |
| PCSrc       | The PC is replaced by the output of the adder that computes the value of PC + 4. | The PC is replaced by the output of the adder that computes the branch target.                          |
| MemRead     | None.  | Data memory contents designated by the address input are put on the Read data output.                   |
| MemWrite    | None.  | Data memory contents designated by the address input are replaced by the value on the Write data input. |
| MemtoReg    | The value fed to the register Write data input comes from the ALU.               | The value fed to the register Write data input comes from the data memory.                              |



# Practice: Fill in this table!

| Inst. | ALUSrc | Memto Reg | Reg Write | Mem Read | Mem Write | Branch |
|-------|--------|-----------|-----------|----------|-----------|--------|
| R-fmt |        |           |           |          |           |        |
| lw    |        |           |           |          |           |        |
| sw    |        |           |           |          |           |        |
| beq   |        |           |           |          |           |        |



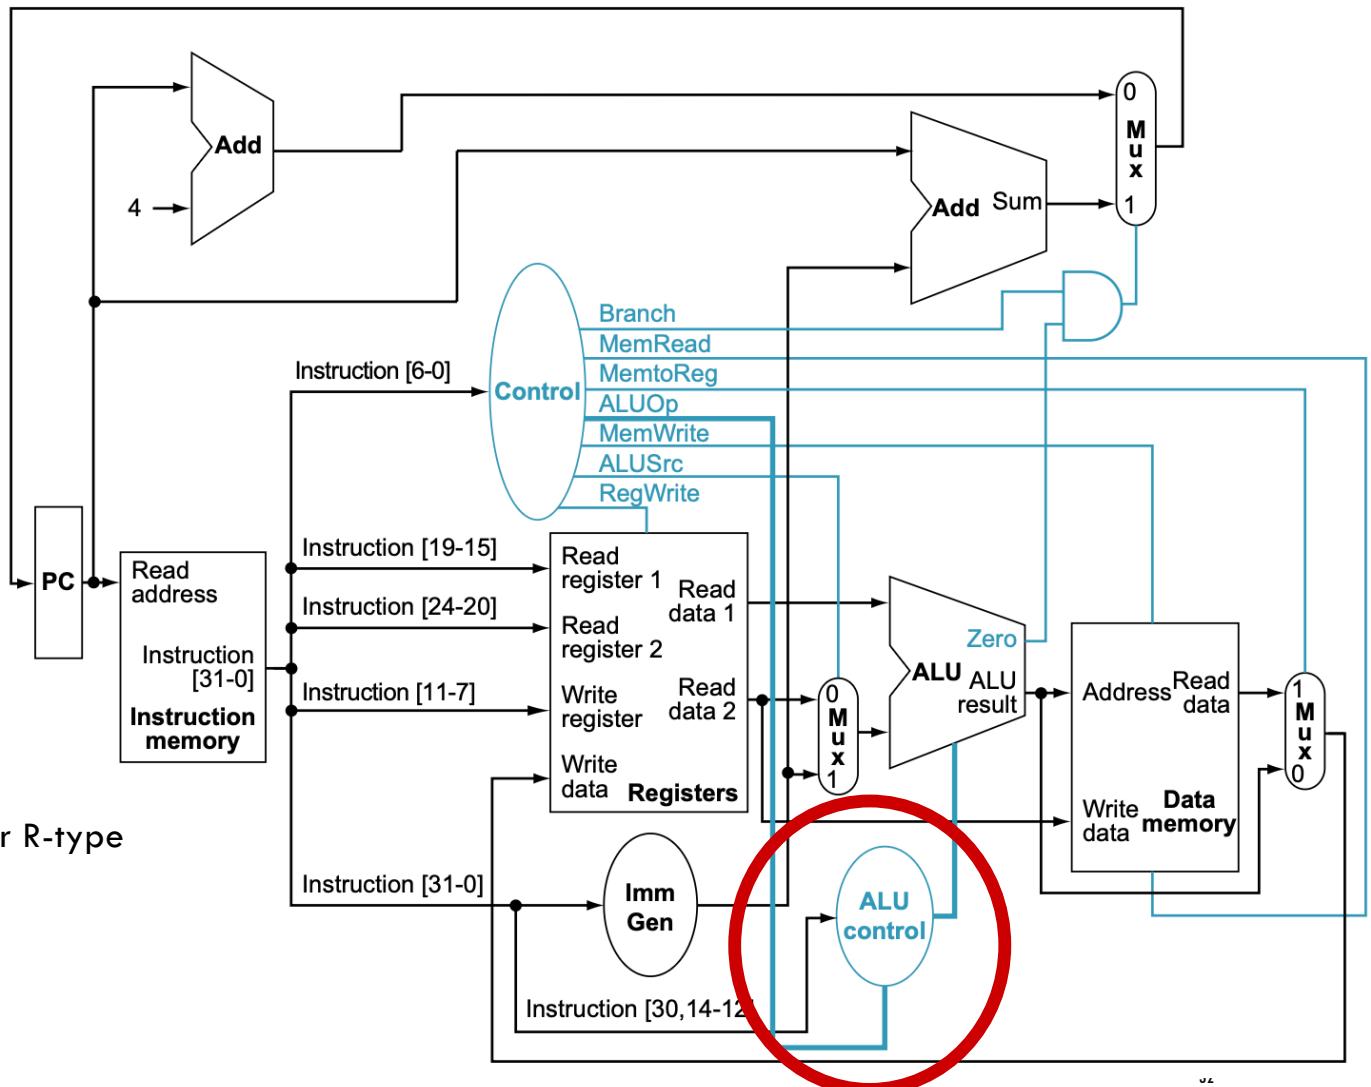
# What is going on here?

| ALU control lines | Function |
|-------------------|----------|
| 0000              | AND      |
| 0001              | OR       |
| 0010              | add      |
| 0110              | subtract |

ALUOp should have ALU control chose:

- **add** for lw/sw
- **subtract** for beq
- **AND/OR/add/subtract** operation for R-type instructions based on its funct fields

[at least needs two bits]

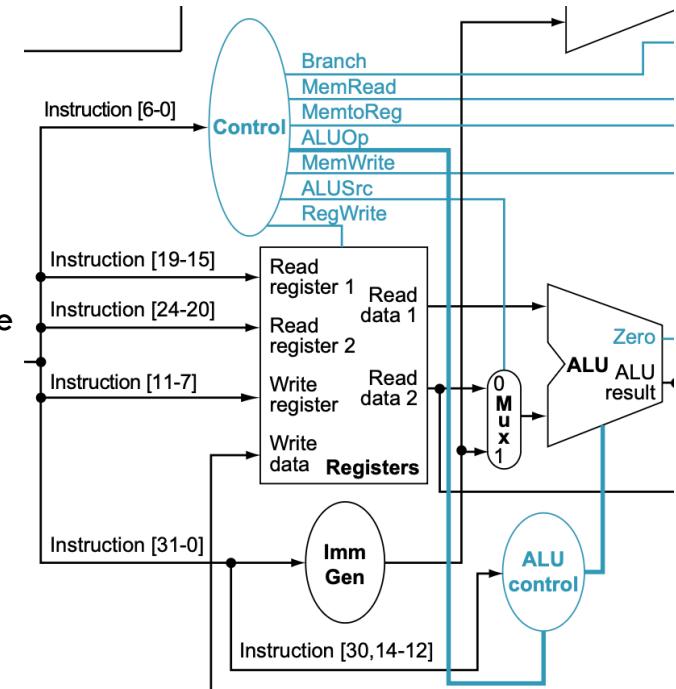


# This is what is going on!

From Before: Only a function of the opcode:

| Instruction | ALUSrc | Memto-Reg | Reg-Write | Mem-Read | Mem-Write | Branch | ALUOp1 | ALUOp0 |
|-------------|--------|-----------|-----------|----------|-----------|--------|--------|--------|
| R-format    | 0      | 0         | 1         | 0        | 0         | 0      | 1      | 0      |
| lw          | 1      | 1         | 1         | 1        | 0         | 0      | 0      | 0      |
| sw          | 1      | X         | 0         | 0        | 1         | 0      | 0      | 0      |
| beq         | 0      | X         | 0         | 0        | 0         | 1      | 0      | 1      |

- add for lw/sw
- subtract for beq
- AND/OR/add/subtract for R-type



| Instruction<br>opcode | ALUOp | Operation       | Funct7<br>field | Funct3<br>field | Desired<br>ALU action | ALU control<br>input |
|-----------------------|-------|-----------------|-----------------|-----------------|-----------------------|----------------------|
| lw                    | 00    | load word       | XXXXXXXX        | XXX             | add                   | 0010                 |
| sw                    | 00    | store word      | XXXXXXXX        | XXX             | add                   | 0010                 |
| beq                   | 01    | branch if equal | XXXXXXXX        | XXX             | subtract              | 0110                 |
| R-type                | 10    | add             | 0000000         | 000             | add                   | 0010                 |
| R-type                | 10    | sub             | 0100000         | 000             | subtract              | 0110                 |
| R-type                | 10    | and             | 0000000         | 111             | AND                   | 0000                 |
| R-type                | 10    | or              | 0000000         | 110             | OR                    | 0001                 |

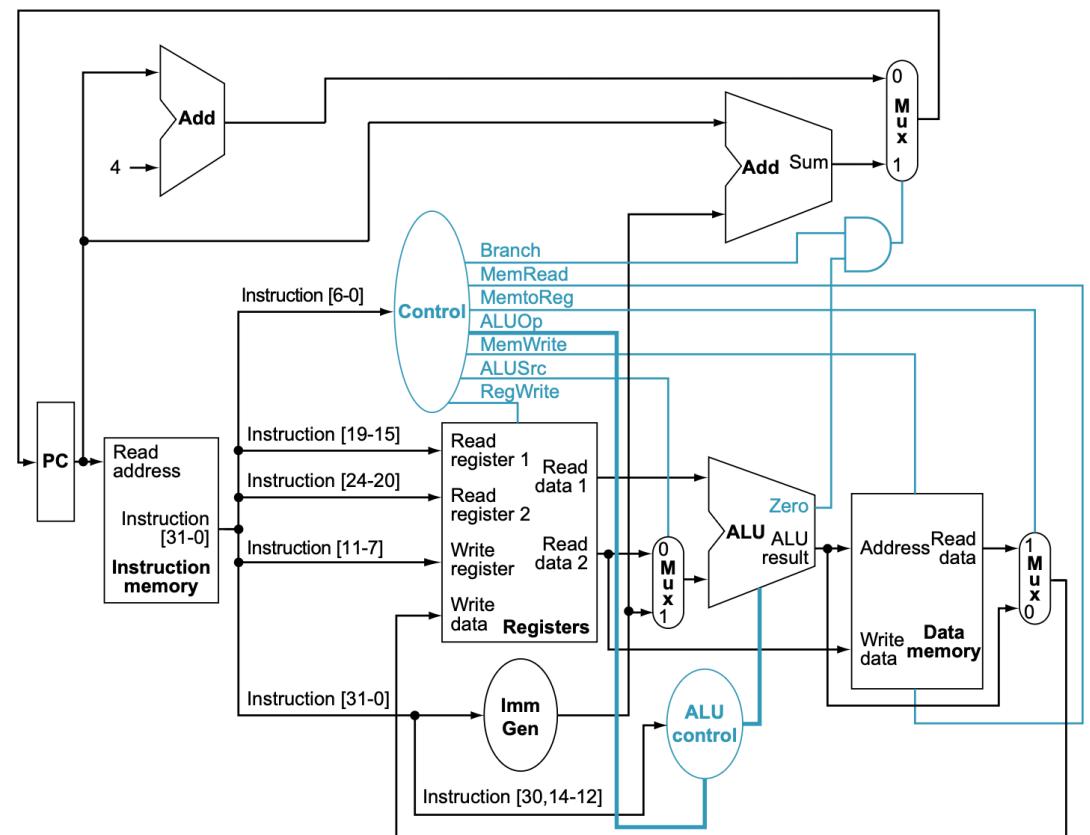
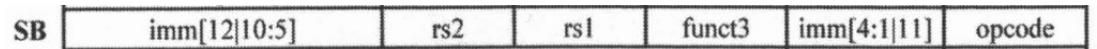
| ALU control lines | Function |
|-------------------|----------|
| 0000              | AND      |
| 0001              | OR       |
| 0010              | add      |
| 0110              | subtract |



# Control Signals: Branch and PCSrc Mux

- PCSrc = Branch (from Control) & Zero (from ALU)
  - True:  $PC = \text{SignExt}(\text{imm}[12:1]) \ll 1 + PC$
  - False:  $PC = PC + 4$
- PC=PC+{imm,1b'0}
- Why?

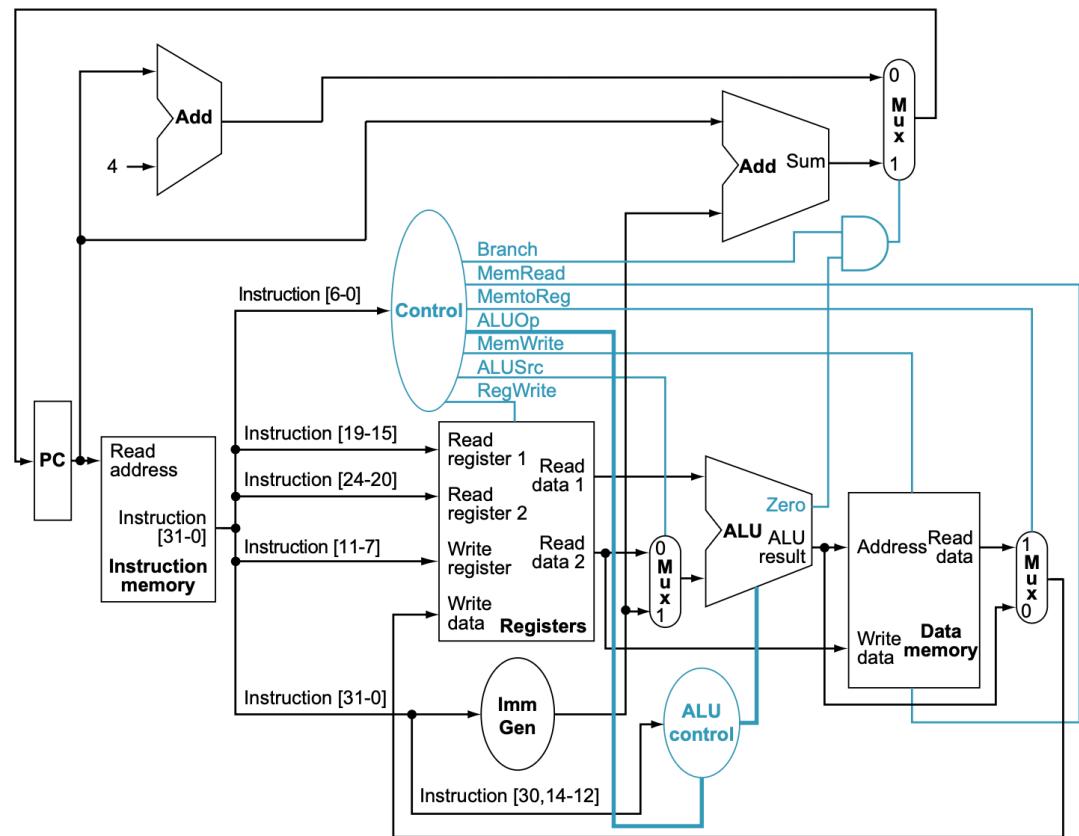
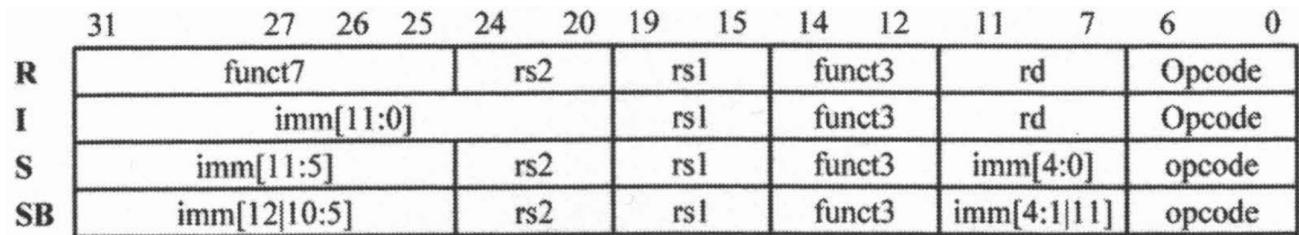
| Instruction | Branch |
|-------------|--------|
| R-format    | 0      |
| lw          | 0      |
| sw          | 0      |
| beq         | 1      |



# Control Signals: ALUSrc

- True:  
ImmGen (lw/I-Format and sw/S-Format)
- False:  
ReadData2 (R-Format, beq/SB-Format)

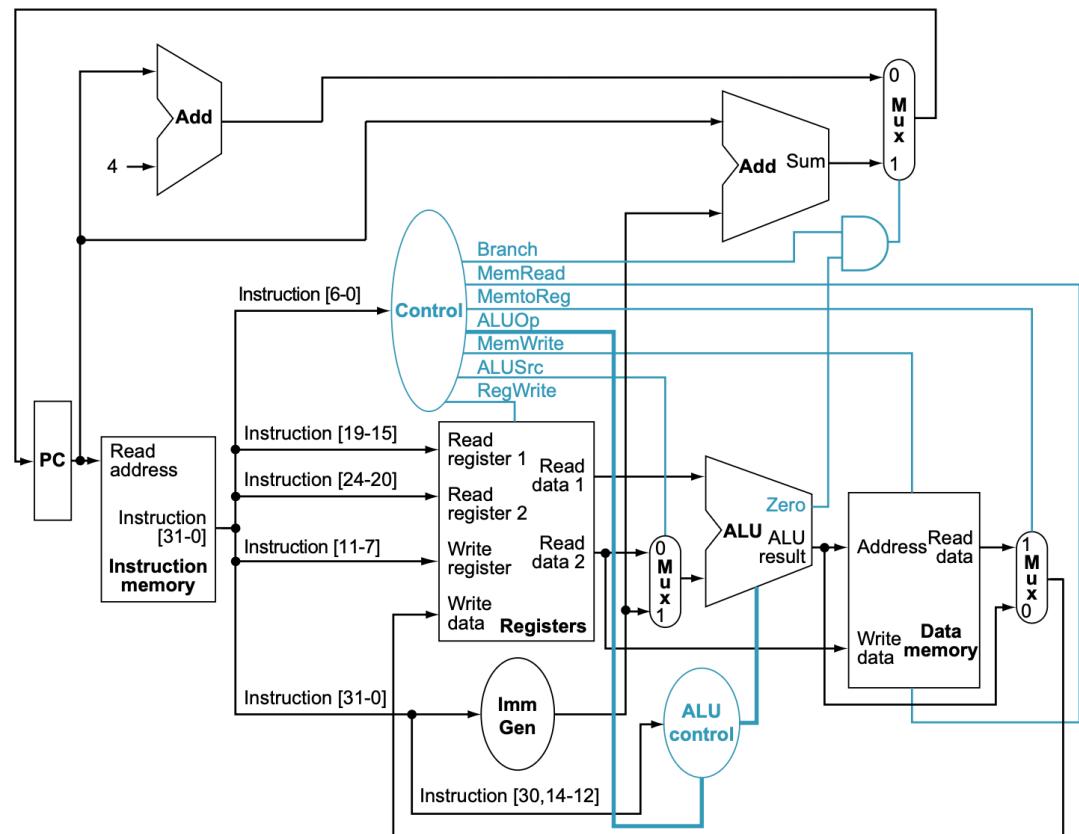
| Instruction | ALUSrc |
|-------------|--------|
| R-format    | 0      |
| lw          | 1      |
| sw          | 1      |
| beq         | 0      |



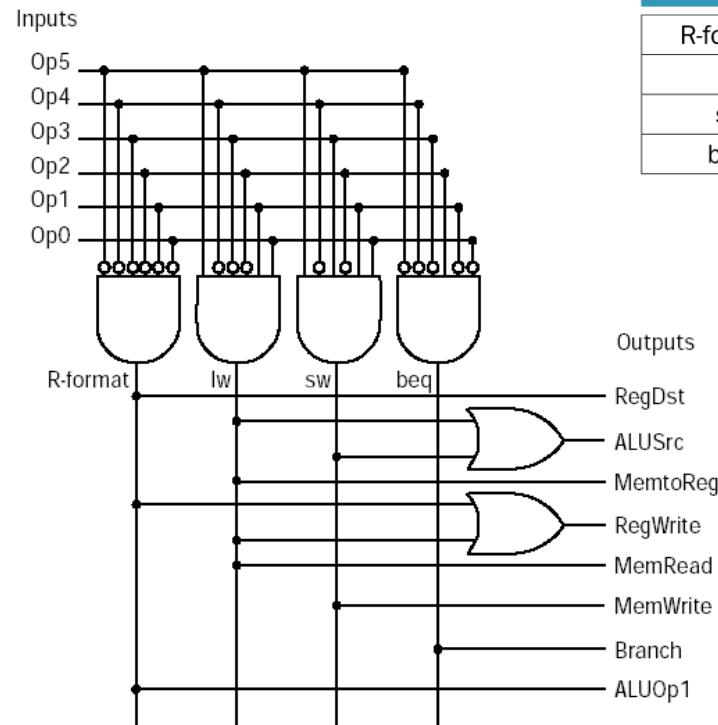
# Control Signals: MemToReg

- True:  
 $\text{RegisterWriteData} = \text{MemReadData}$
- False:  
 $\text{RegisterWriteData} = \text{ALUResult}$

| Instruction | Memto-Reg |
|-------------|-----------|
| R-format    | 0         |
| lw          | 1         |
| sw          | X         |
| beq         | X         |



# Control Design



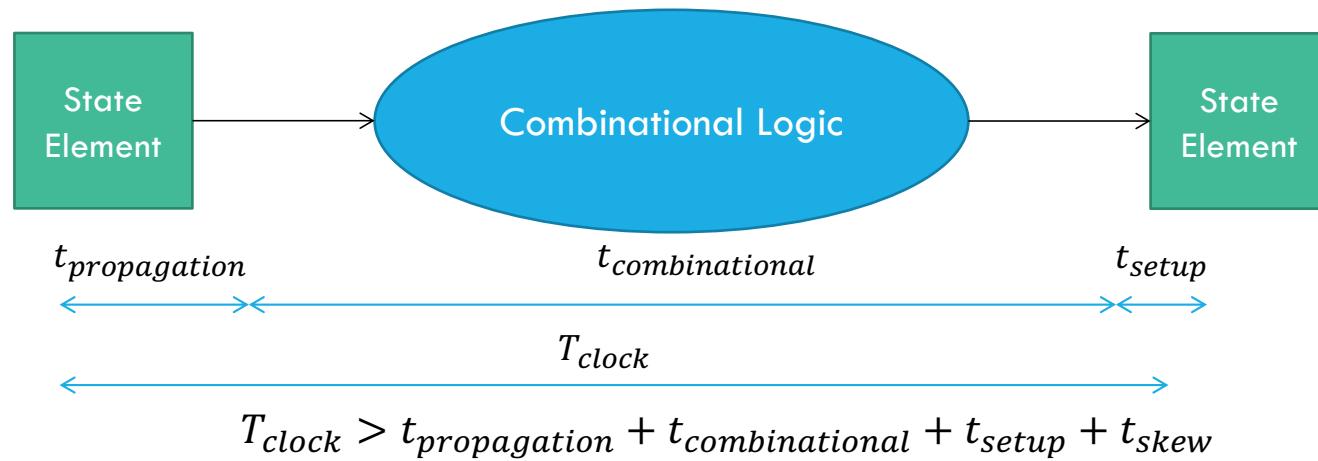
| Instruction | ALUSrc | Memto-Reg | Reg-Write | Mem-Read | Mem-Write | Branch | ALUOp1 | ALUOp0 |
|-------------|--------|-----------|-----------|----------|-----------|--------|--------|--------|
| R-format    | 0      | 0         | 1         | 0        | 0         | 0      | 1      | 0      |
| lw          | 1      | 1         | 1         | 1        | 0         | 0      | 0      | 0      |
| sw          | 1      | X         | 0         | 0        | 1         | 0      | 0      | 0      |
| beq         | 0      | X         | 0         | 0        | 0         | 1      | 0      | 1      |

Memory Reference: lw, sw  
 Arithmetic/Logical: add, sub, and, or, slt  
 Control flow: beq

Truth Table (above) -> SoP -> Circuit

From MIPS, how can you tell? Works with RISC-V!

# Single-Cycle Implementation



Cycle time determined by the length of the longest path.



What is longest state element to state element path?

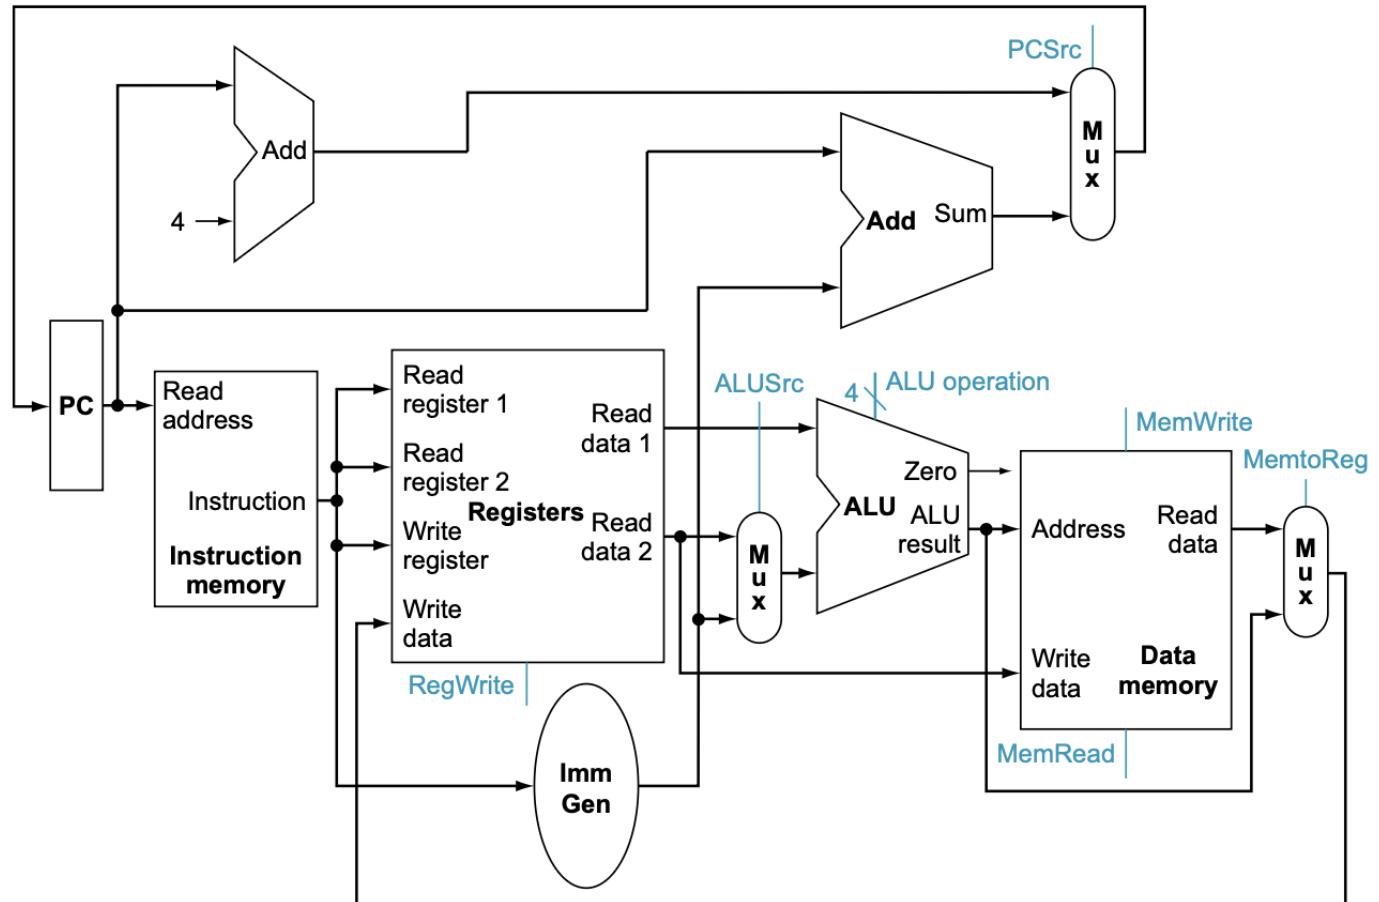
# Single-Cycle Implementation

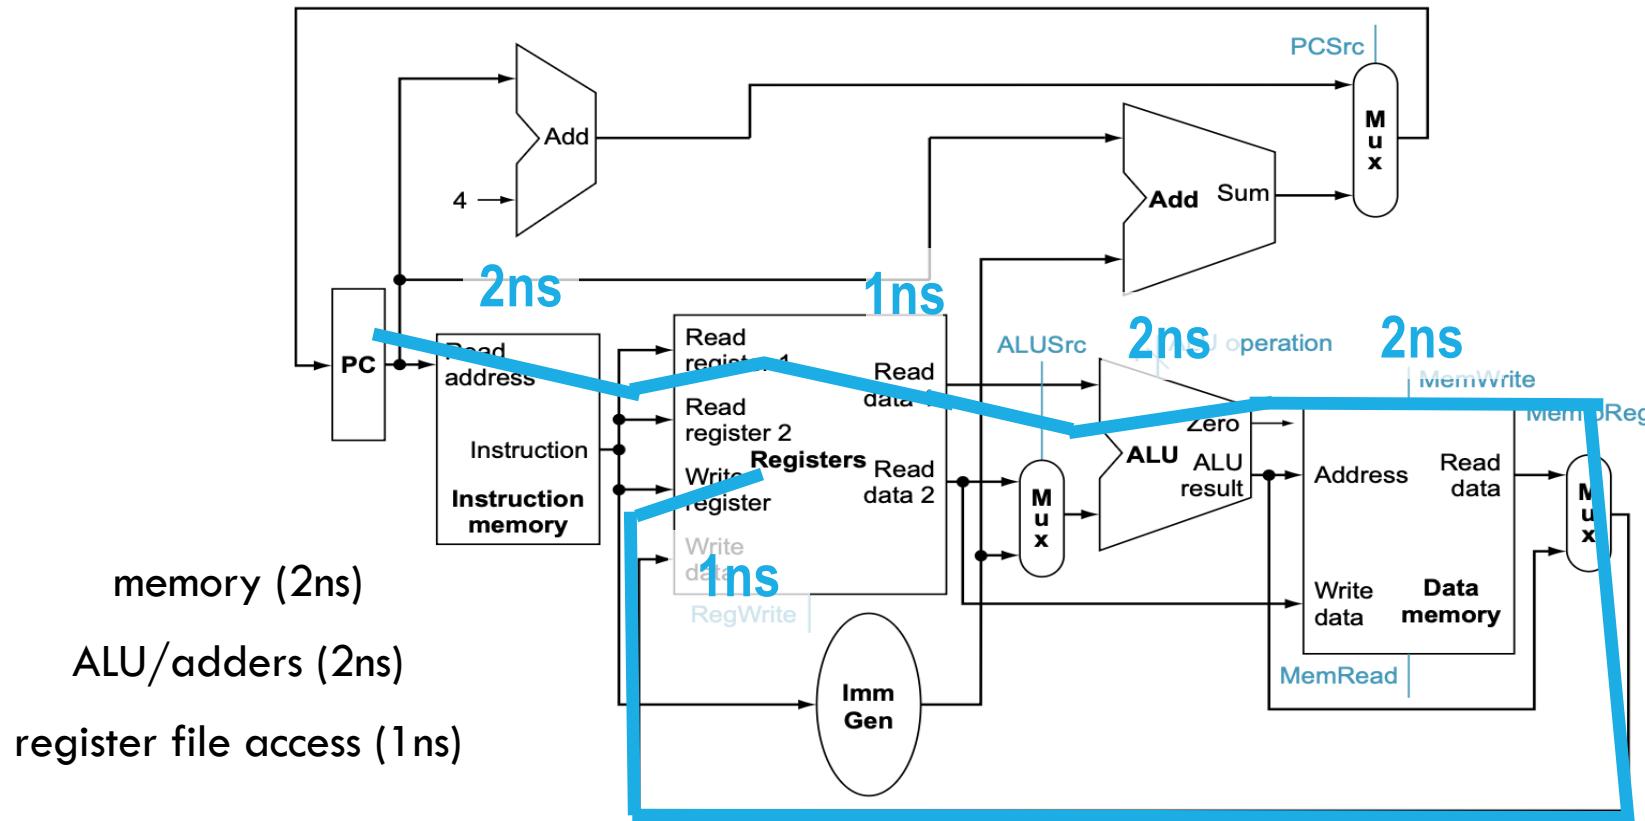
Calculate cycle time  
assuming negligible  
delays except:

memory (2ns)

ALU/adders (2ns)

register file access (1ns)





Load Word is longest running instruction  
(prove by computing time for sw, beq, R-type)

$$2\text{ns} + 1\text{ns} + 2\text{ns} + 2\text{ns} + 1\text{ns} = \mathbf{8\text{ns}}$$

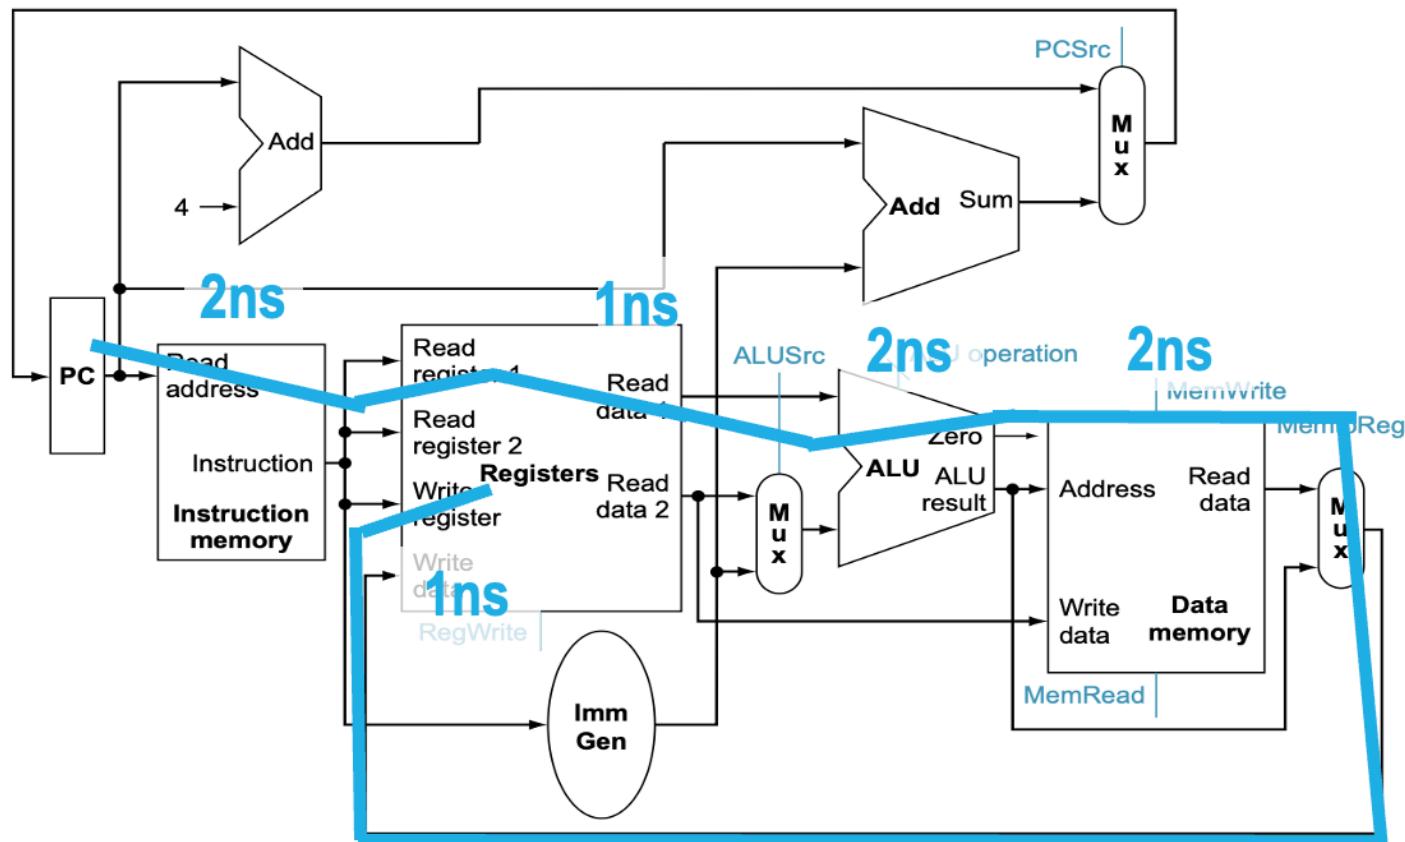
# Magic Moment!

At this point, you can design a processor!

1. Analyze instruction set for datapath requirements
2. Select set of datapath components
3. Assemble datapath meeting the requirements
4. Analyze implementation of each instruction to determine setting of control signals
5. Assemble the control logic
6. Set clock rate based on longest path



# What's wrong with the single-cycle implementation?



# What's wrong with the single-cycle implementation?

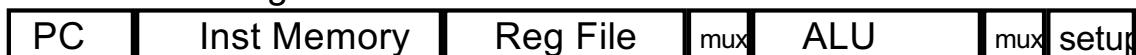
memory (2ns), ALU/adders (2ns), register file access (1ns)

| Inst. | Inst. Mem | Reg. Read | ALU | Data Mem | Reg. Write | Total |
|-------|-----------|-----------|-----|----------|------------|-------|
| ALU   | 2         | 1         | 2   |          | 1          | 6     |
| lw    | 2         | 1         | 2   | 2        | 1          | 8     |
| sw    | 2         | 1         | 2   | 2        |            | 7     |
| beq   | 2         | 1         | 2   |          |            | 5     |

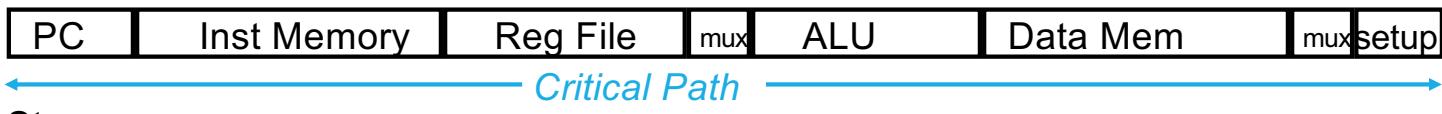


# What's wrong with the single-cycle implementation?

Arithmetic & Logical



Load



Store



Branch



- Long cycle time, much of it wasted
- All instructions take as much time as the slowest
- Real memory is not so nice as our idealized memory
  - (cannot always get the job done in fixed amount of time)



# How bad is it?

Assume: 100 instructions executed

- 25% of instructions are loads (8ns),
- 10% of instructions are stores (7ns),
- 45% of instructions are adds (6ns), and
- 20% of instructions are branches (5ns).

Single-cycle execution:

$$100 * 8\text{ns} = 800 \text{ ns}$$

Optimal execution:

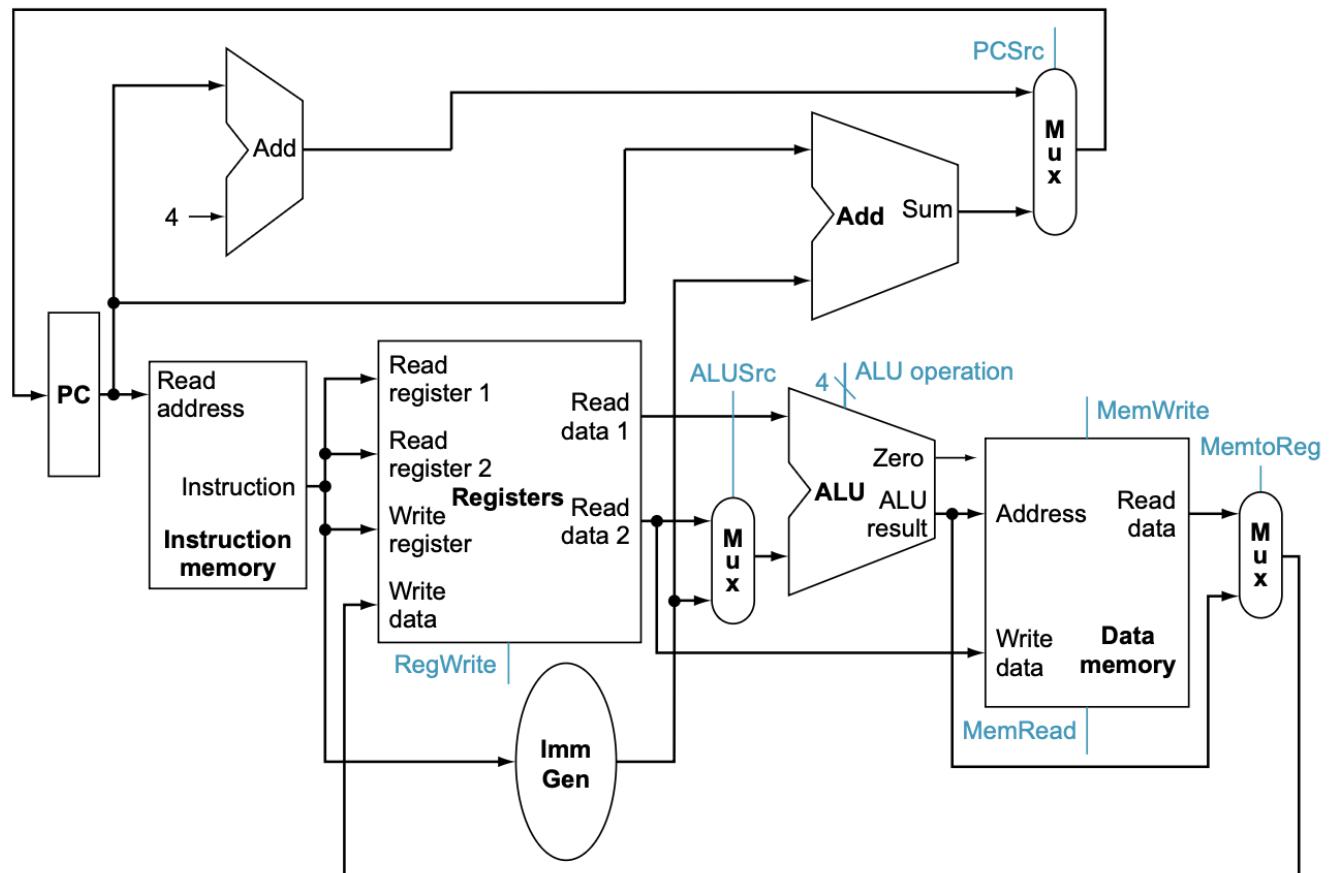
$$25*8\text{ns} + 10*7\text{ns} + 45*6\text{ns} + 20*5\text{ns} = 640 \text{ ns}$$

$$\text{Speedup} = 800/640 = 1.25$$



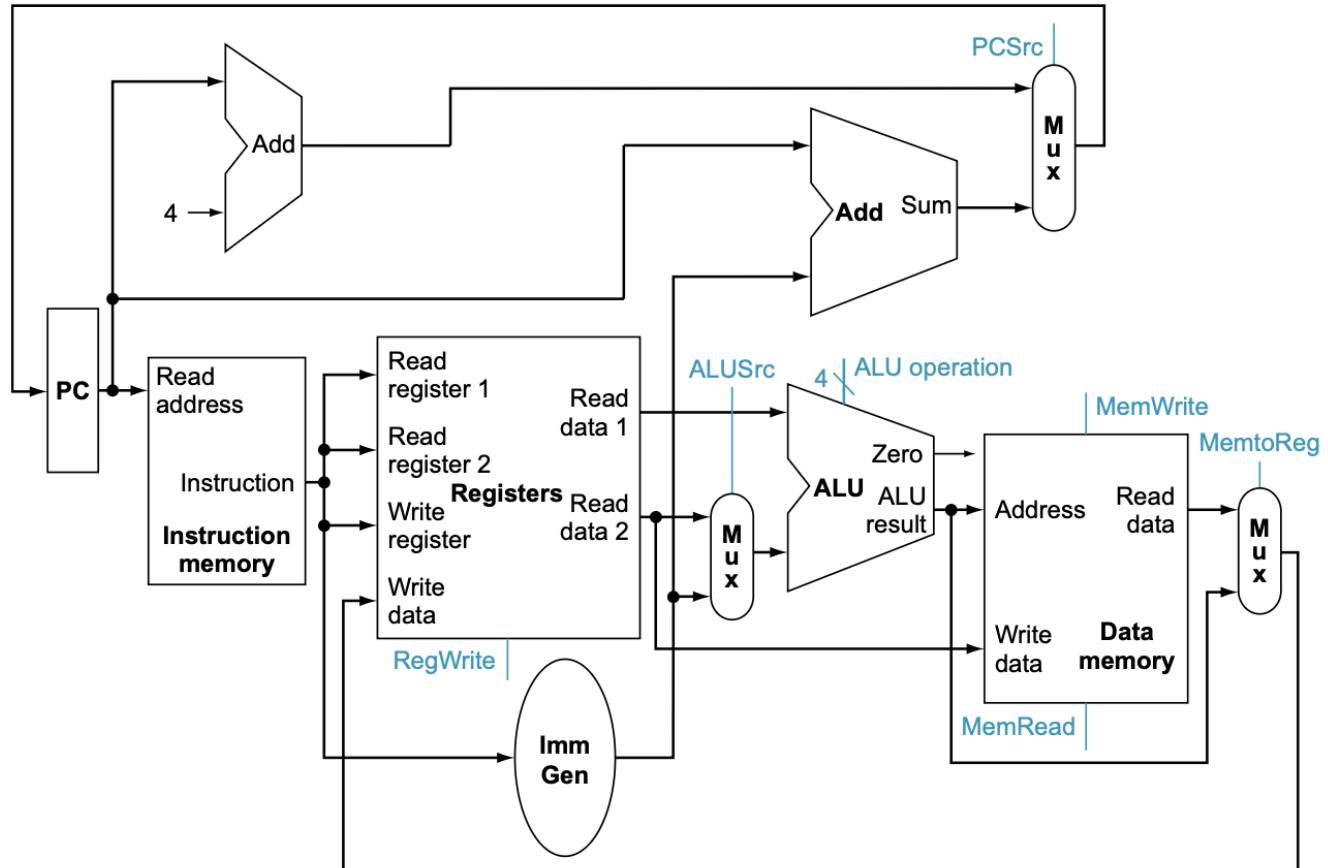
# Other Issues

- Instruction and Data Memory are the same
- Some units hold value long after job is complete
  - Could reuse ALU for example
  - Underutilized resources (wasteful of area/power)



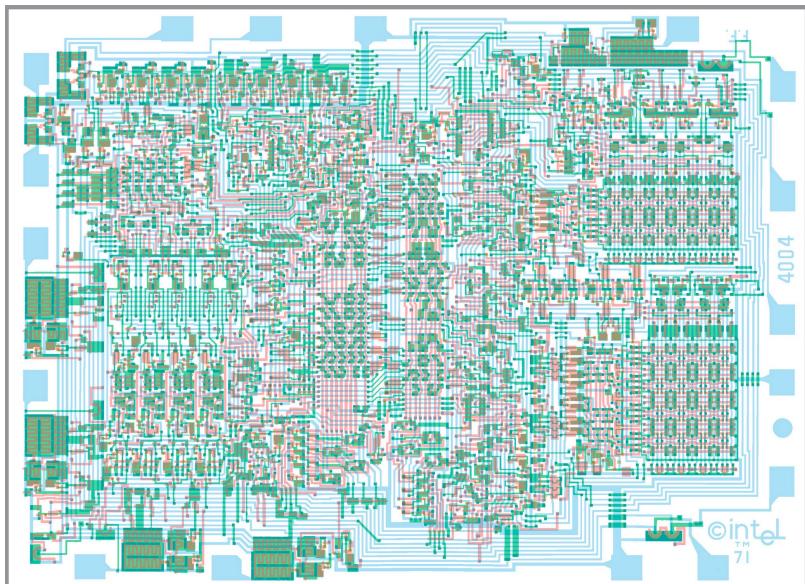
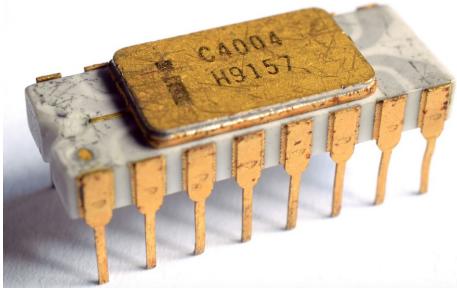
# Other Issues

- Floating point and other VERY LONG instructions



# Welcome to 1971!!!

## 1949 was nice, but...



48

### The Intel 4004, 1971:

- **First microprocessor  
(stored-program, electronic, general-purpose)**
- 4-bit CPU with 2,250 transistors, 2 designers
- 10.8  $\mu$ Second instruction cycle
- <http://www.intel4004.com/>   <http://e4004.szyc.org/>
- Full design recently released by Intel

Intel introduces an integrated CPU chip with a 4-bit parallel serial interface. It requires no accumulators and a push-down stack on one chip. It's one of a family of four microprocessors—the first complete computer system—the first system to bring you the power and economy of a mainframe computer at low cost in as few as two dual-in-line packages.

MCS-4 systems provide complete computing and control functions for business, scientific, billing machines, measuring systems, numeric control systems and other specialized parts.

The heart of any MCS-4 system is a Type 4004 CPU, which includes a powerful set of 45 instructions. Adding one or more of the other three chips—Type 4002 ROM and data tables gives you a fully functioning micro-computer system. Type 4003 provides 16K of 8-bit RAM for read/write memory and Type 4002 registers to handle memory operations.

Using no circuitry other than ICs from the family of four, you can build a system with 16K bytes of ROM storage and 5120 bits of RAM storage. When you consider that the 4004 chip contains a micro-programmed ROM, Intel's invisible and re-programmable ROM, Type 1000, and the 4003 ROM, it's clear that the micro-programmed ROM.

MCS-4 offers performance even in switches, logic boards, displays, teletypewriters, printers, readers, A-D converters and other peripheral equipment.

The MCS-4 family is now in stock at Intel's Santa Clara headquarters and at our marketing headquarters in Foster City, California, and at our international representative for technical interests and sales, B. F. Goodrich Electronics Division, 1000 Lincoln, 1000 Lincoln, Bechtel Parkway, Mountain View, No. 4-2000, Mountain View, California 94031. Phone 415-960-2747. Intel Corporation also manufactures microprocessor memory devices and memory systems at 3050 Bowery Avenue, San Jose, Calif. 95131. Phone 408-265-7001.

**intel delivers.**

The first advertisement for a microprocessor in Electronic News in November 1971. [1]

**FIN**



Please fill out ICQ Form here!



<https://tinyurl.com/Fa25ICQ>

